

# CIS 4930-001: INTRODUCTION TO AUGMENTED AND VIRTUAL REALITY

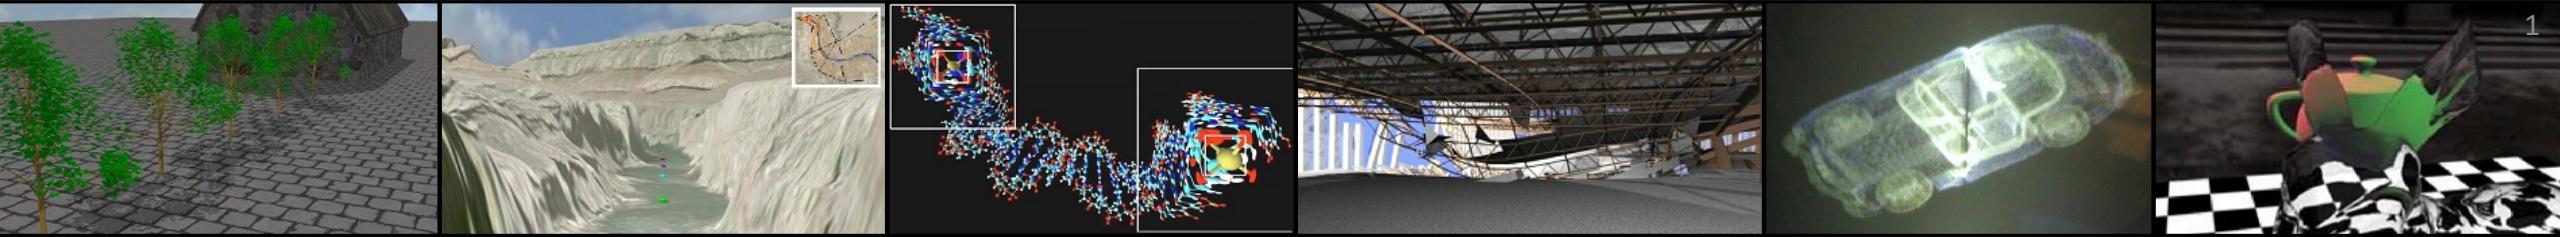
---



## Scene Graph and Acceleration Strategies

Paul Rosen  
Assistant Professor  
University of South Florida

Some slides from: Anders Backman, Mark Billinghurst, Doug Bowman, David Johnson, Gun Lee,  
Ivan Poupyrev, Bruce Thomas, Geb Thomas, Anna Yershova, Stefanie Zollman



# MOTIVATION

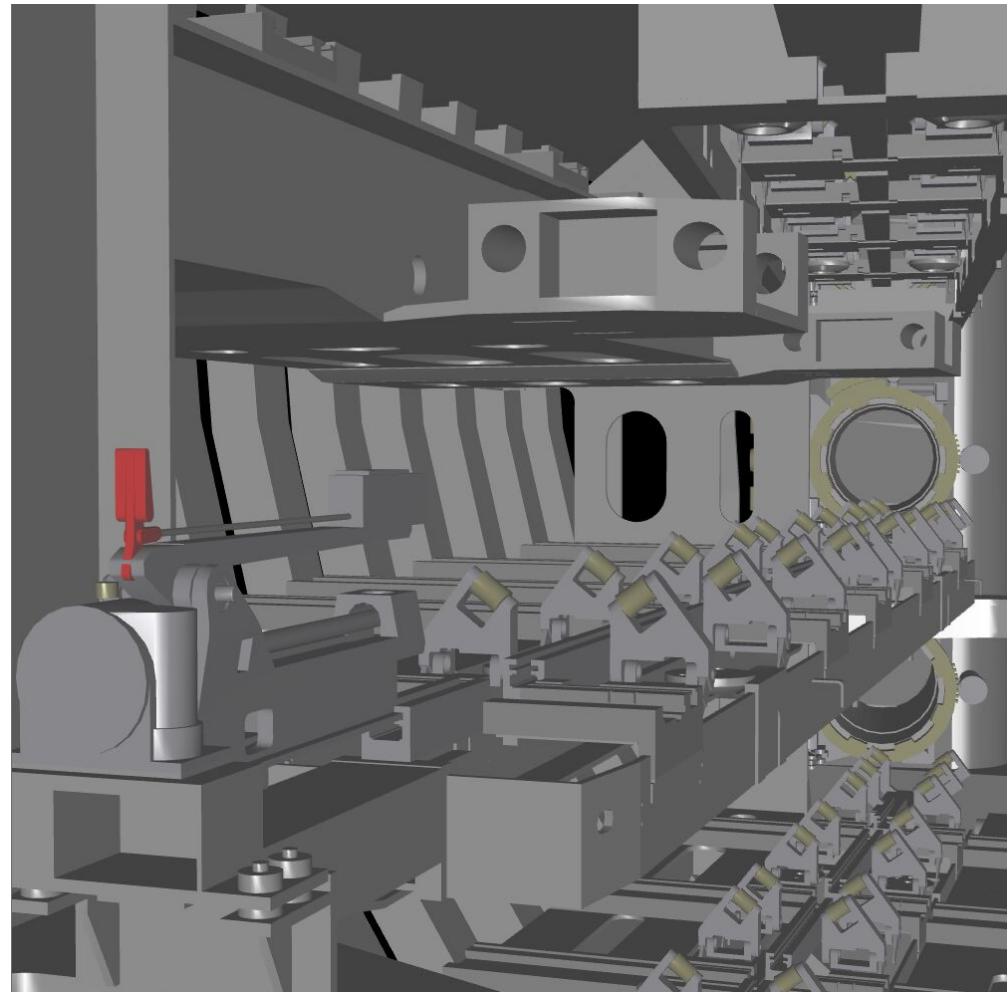
VR scenes need complexity

- Immersion
- People don't bother with VR for easy problems
  - lots of geometry

How do we organize objects?

How do we maintain speed?

- Wait for the graphics cards?



## LIMITATIONS OF IMMEDIATE MODE GRAPHICS

When we define a geometric object in an application

Object is passed through the pipeline

To redraw the object, either changed or the same, we must  
re-execute the code

At the end of the day, we manage objects directly



# THE SCENE GRAPH

Build a structure that allows:

- More intuitive/organized modeling
- Global analysis of the model
  - To generate OpenGL calls in an efficient way
- CPU-level preprocessing to avoid overloading graphics pipeline
  - Object culling



# SCENE GRAPHS

Thinking objects ... not vertices

- Scene Graph: Directed Acyclic Graph (DAG)
  - Represents object-based hierarchy of geometry
  - Suits a bottom-up design of objects & scenes
  - Common for most graphic API's

Thinking content ... not rendering process

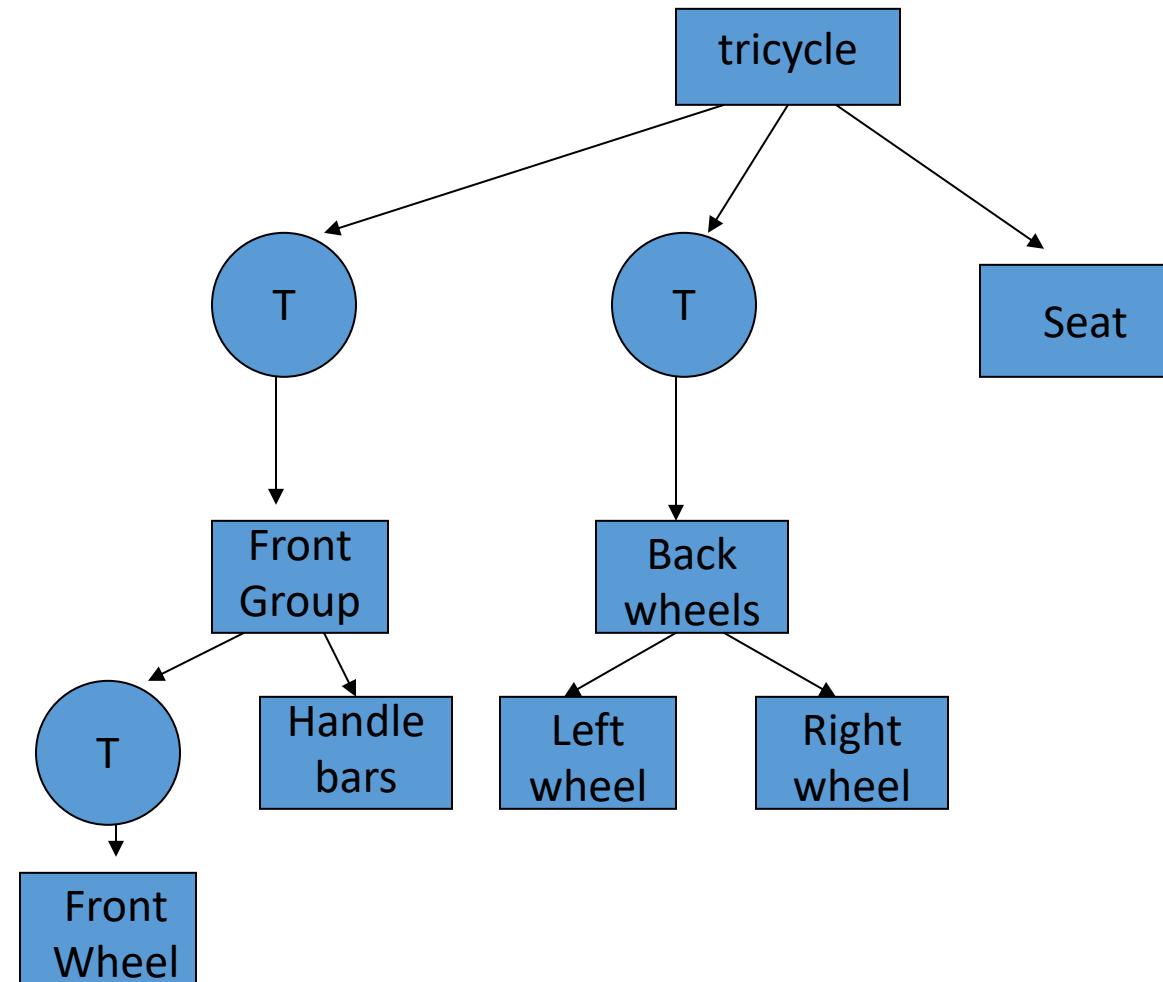
- A scene is a collection of nodes
  - Nodes are: groups, lights, geometry, transforms, sounds, rendering states, etc...
  - Nodes hold pointers to children

Provides the hierarchical framework for grouping objects together, spatially

- Behavior can be built into the scene-graph: collision detection, animations, ...
- Abstracts the hardware (i.e. multipass/multitexture)
- Possible to optimize the rendering
  - Sorting of: material state changes; polygons when rendering transparency; etc.



# EXAMPLES SCENE GRAPH



## EXAMPLE OF SOME SCENE-GRAPH ELEMENTS

### Group Nodes

- Collects a number of nodes in a logical way
- Level of Detail, Switch nodes

### Transformation Nodes

- Moves, rotates, scales its children/siblings

### Light Nodes

- Point light, Spotlight, Directional, ...

### Geometry Nodes

- Usually a leaf node
- Polygons, lines, points, ...
- Material

### Others

- Sound, fog, manipulators (animators), Collision, ...

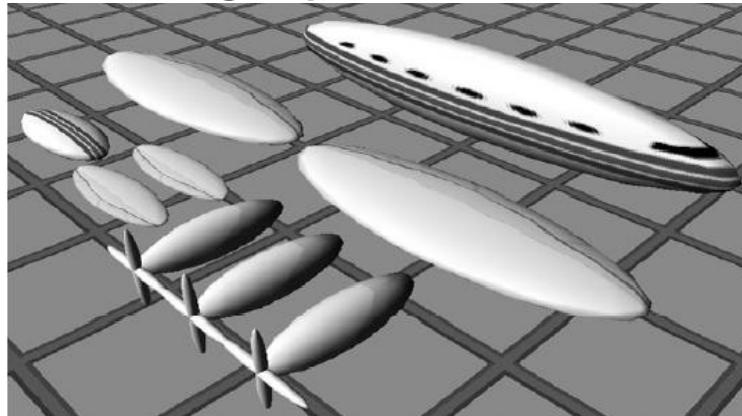
### Visitor/Traversers

- Walks through the scene graph and operate on each node
- RenderVisitor, CullVisitor, ...

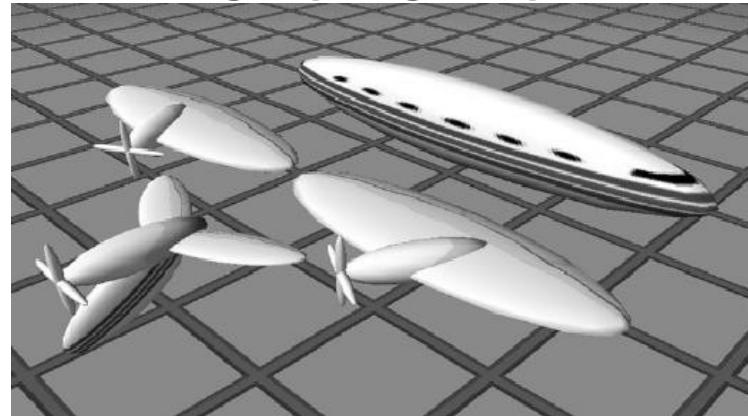


# AN EXAMPLE

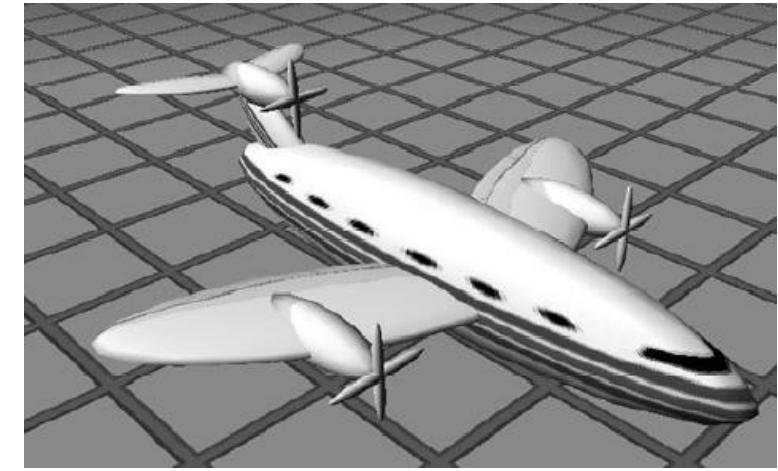
Scene graph items



Scene graph groups



Final Scene



## BENEFITS

Transparent use of underlying hardware

Possibly for the scene graph to optimize the rendering of the content

- State sorting (render objects with same shaders, textures, etc. together)
- Rendering order of transparent objects

Sharing of data

- Objects with the same appearance can share state
- Objects with the same geometry can share geometry (smaller memory footprint)

Possible to serialize (and to store or send over the internet)

Rendering Acceleration!



# ACCELERATION ALGORITHMS

Never enough horse power

Four performance goals

- Higher FPS
- Higher Resolution
- More realism
- Lower power

We will always need to assist the hardware with clever data structures and algorithms



## SPATIAL DATA STRUCTURES

Organize data in some n-dimensional space

We want to accelerate queries such as

- Does the following geometry overlap any other geometry?
- Does this ray intersect anything in the scene?

Using a tree structure

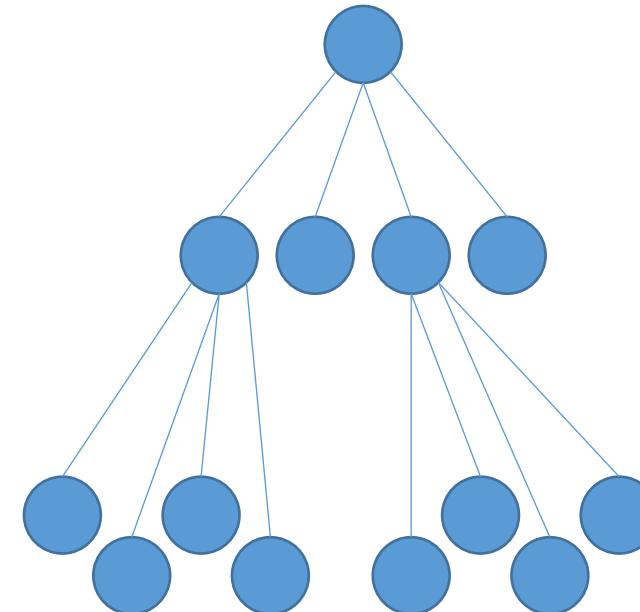
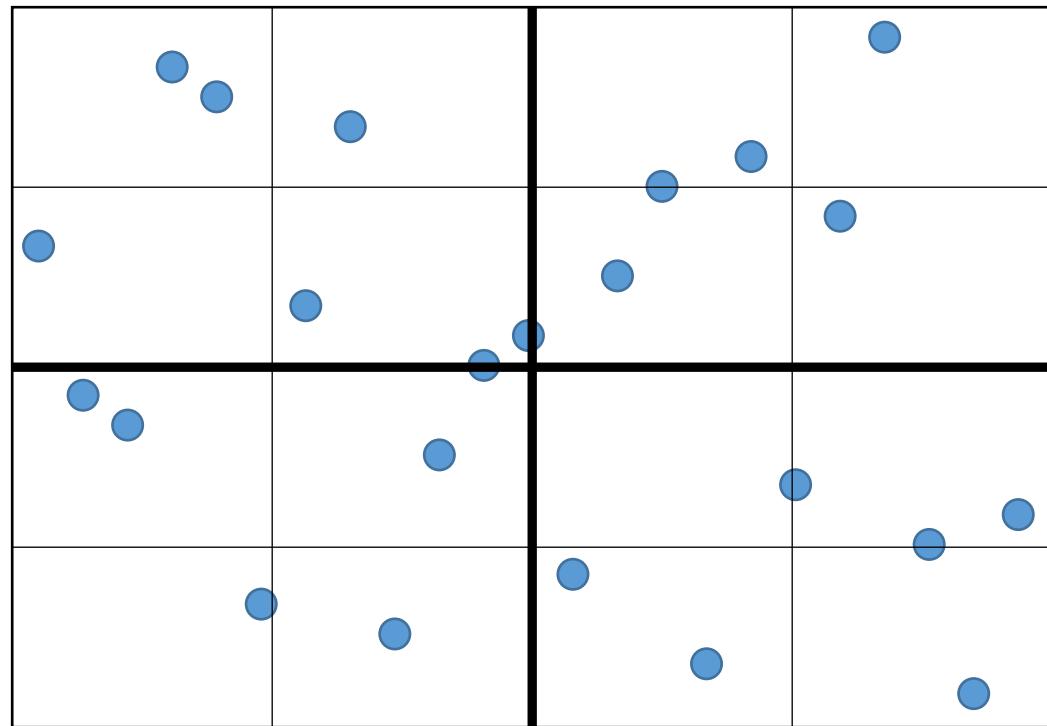
- Construction of a tree can be expensive and should in some cases be done as a preprocess
- Speedups for queries and traversals from  $O(n)$  to  $O(\log n)$
- Three flavors:
  - Binary Space Partition trees (BSP)
  - Quadtrees/Octrees/Kd-trees
  - Bounding Volume Hierarchies (BVH)



# QUADTREES/OCTREES

Recursively divide space in 2D (quadtrees) or 3D (octrees)

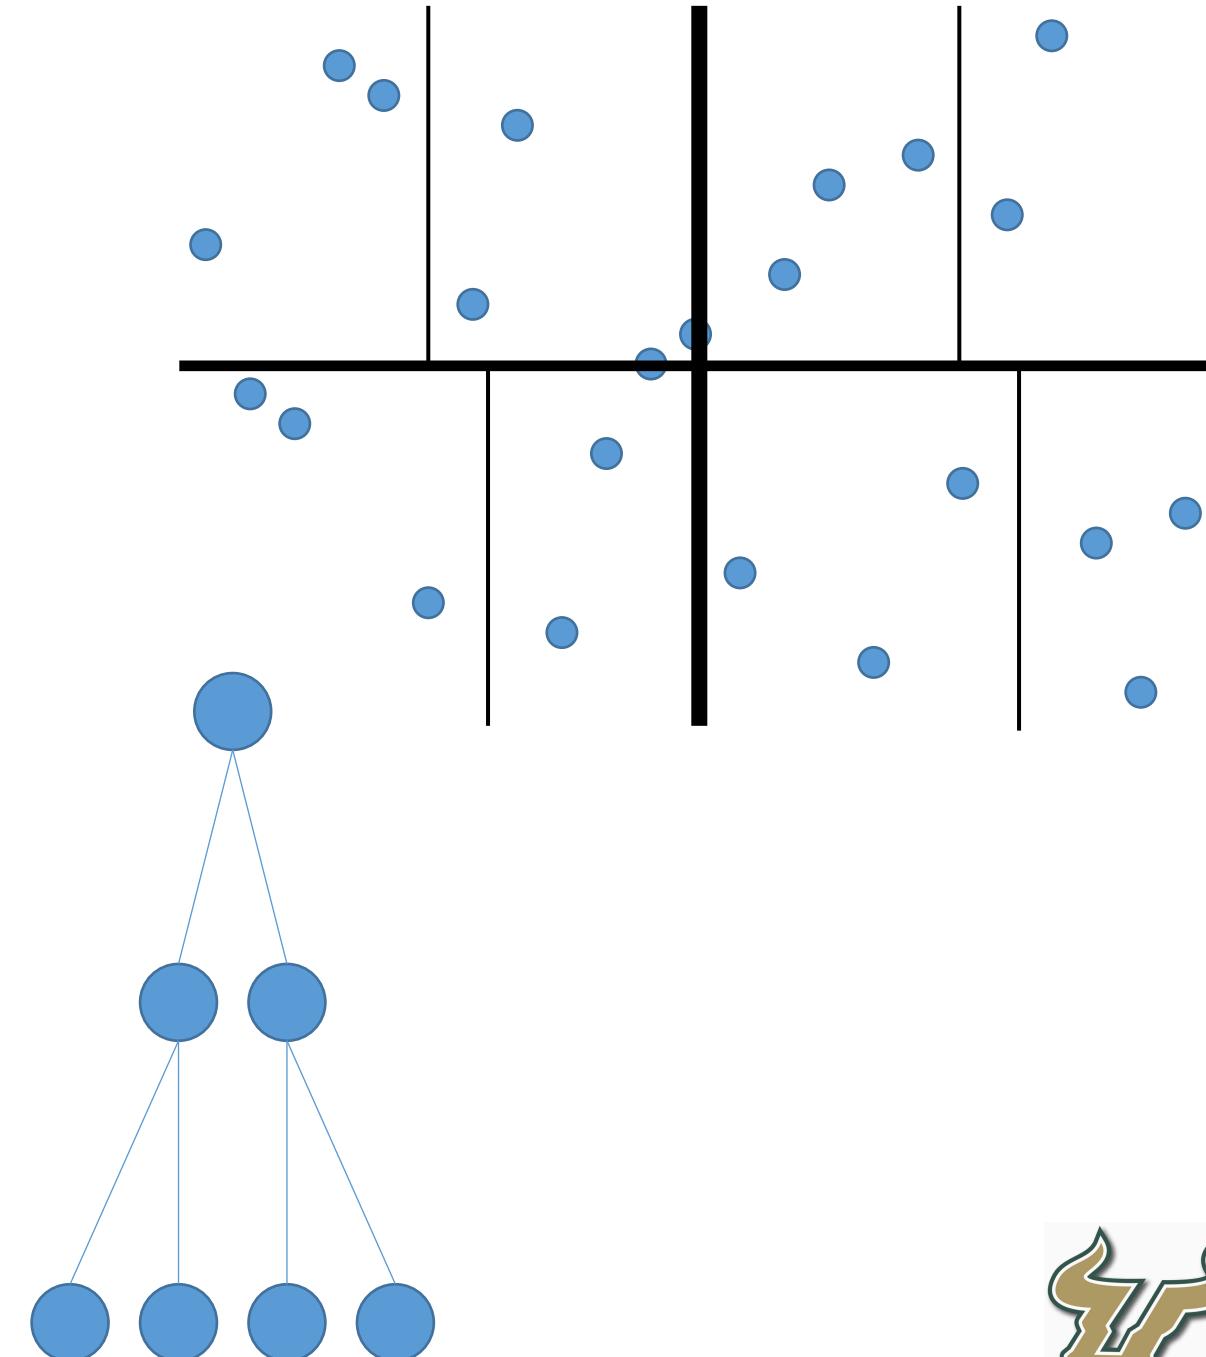
At each tree level objects get placed in the quadrant or octant that is “most appropriate”.



## KD-TREES

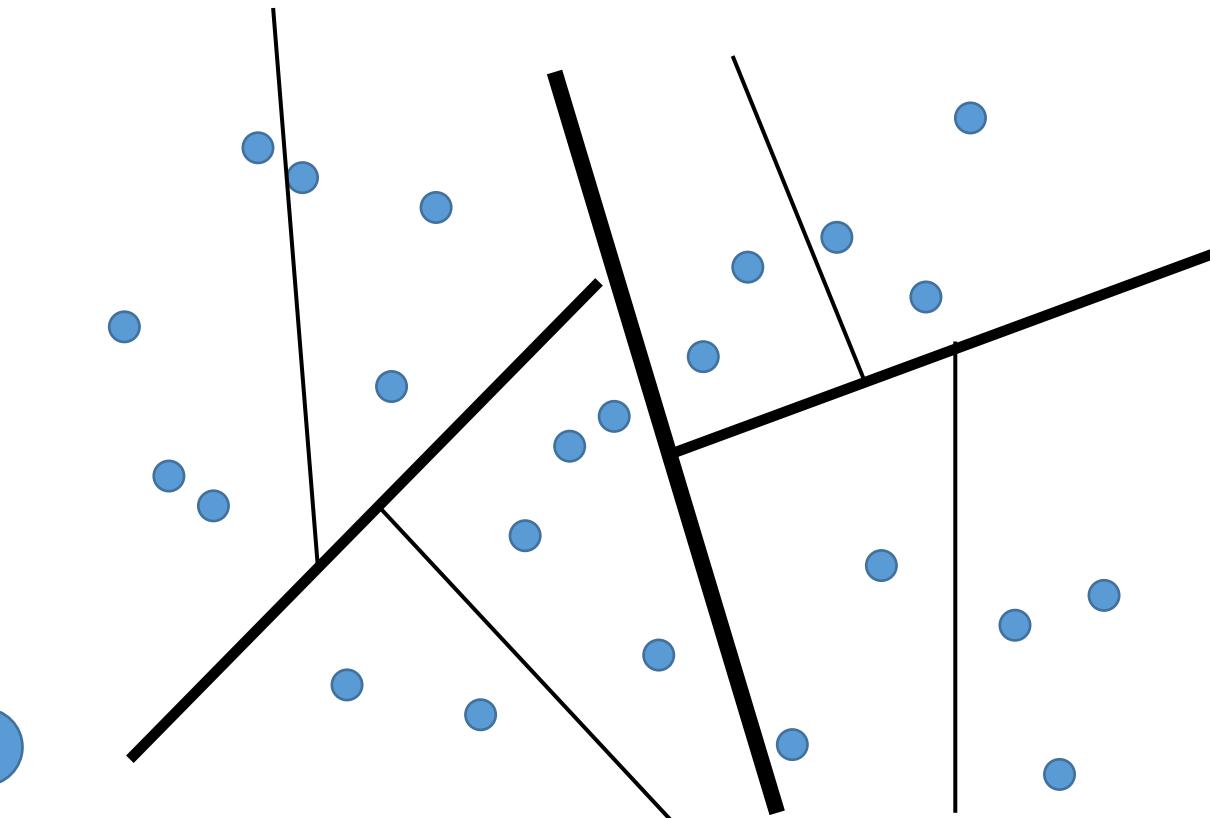
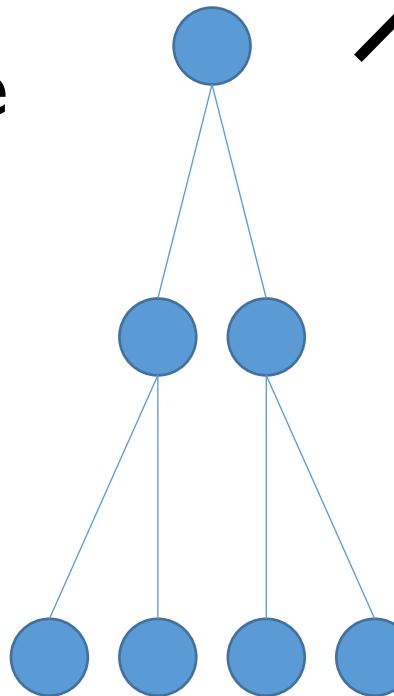
Recursively divide space in 2 parts by selecting the dimension and cut that best divides the objects

At each level objects get placed in the “most appropriate” division.



## BINARY SPACE PARTITION (BSP)

Similar to kd-tree except  
cuts no longer need to be  
axis aligned



## BOUNDING VOLUMES

We want a simpler representation of a geometry

- Not for rendering, but for intersection/collision tests

Specs:

- Should fit enclosed geometry as tight as possible
- Should be simple to use in collision test
- Specs are contradictory



# BOUNDING VOLUME GEOMETRY

## Bounding spheres

- + Simple to create
- + Simple collision detection
- - Could potentially be a bad choice for long thin objects

## Axis Aligned Bounding Box (AABB)

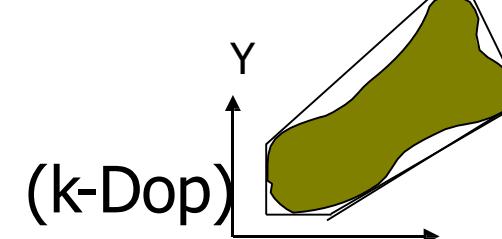
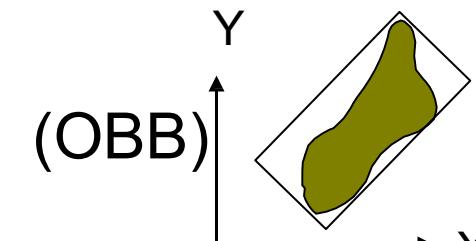
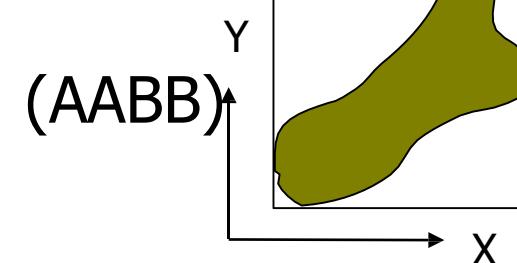
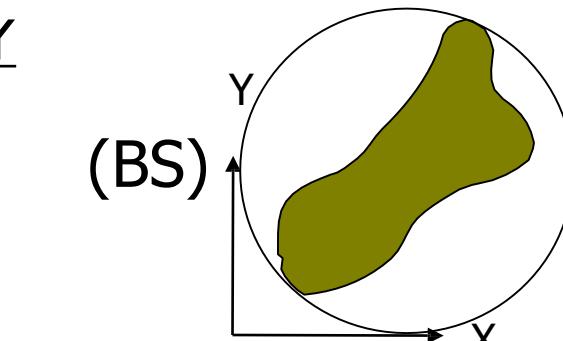
- + Simple to create
- + Simple collision detection
- - A lot of empty space enclosed

## Oriented Bounding Box (OBB)

- + Less empty space enclosed
- - More complicated to create
- - More complicated during collision detection

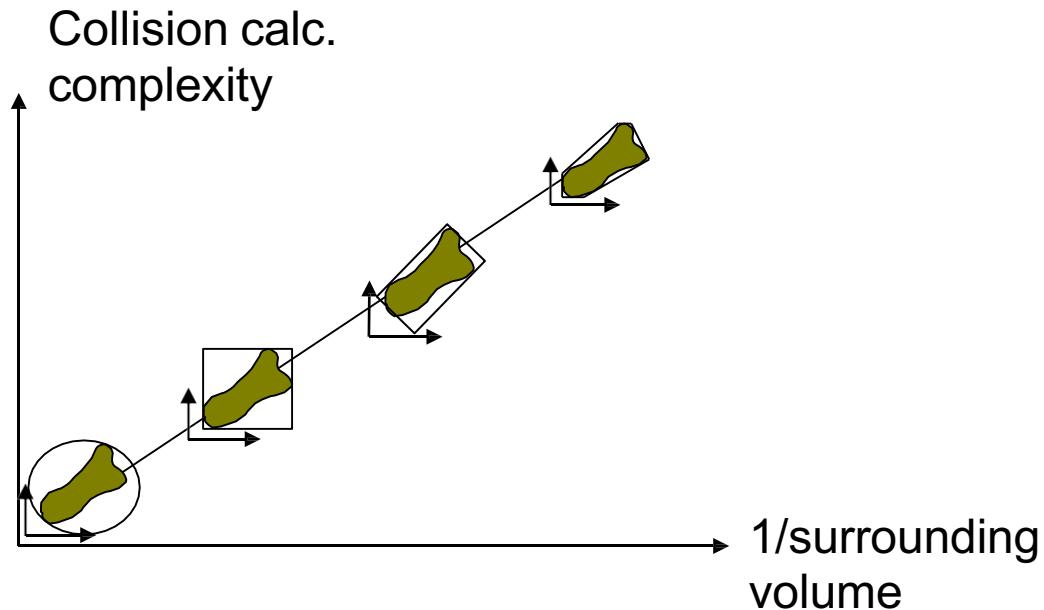
## k-Dop (discrete oriented polytope)

- + Least empty space
- - Most complicated to calculate



# BOUNDING VOLUME GEOMETRY

Less empty space – more complex to create and use



*“The fastest rendered polygon, is the one  
never sent down the rendering pipeline”*

Real-Time Rendering  
Akenine-Möller, Haines, and Hoffman



## SPEEDING THE GRAPHICS PIPELINE

What are some ways the graphics pipeline speeds rendering?

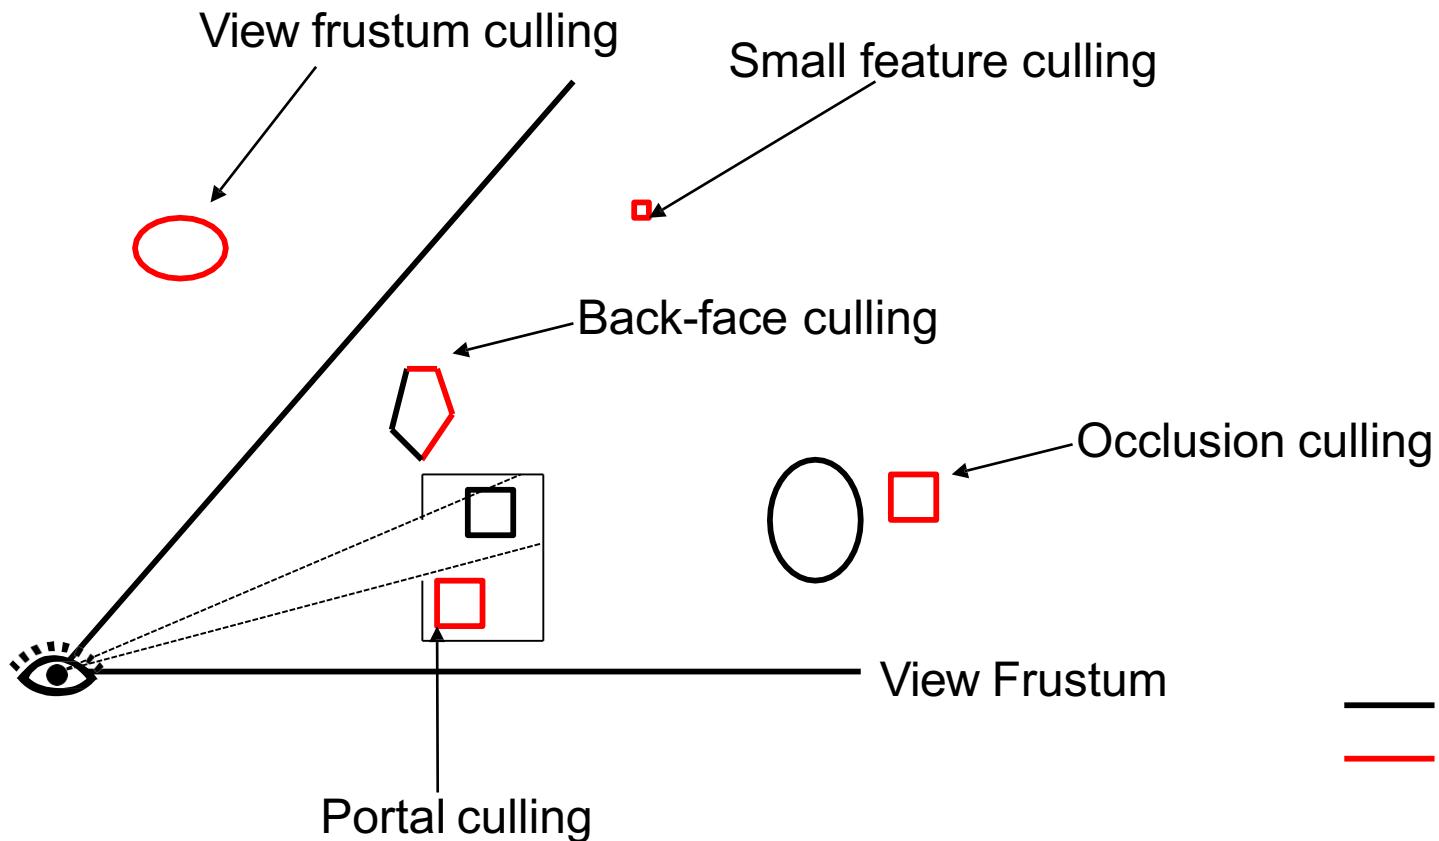
- Off-screen geometry: solved by clipping
- Occluded geometry: solved by Z-buffer
- Clipping/Z-buffering take time linear to the number of primitives/fragments

How can Scene Graphs help?

- Off-screen: **view-frustum culling, portal culling, contribution culling**
- Occluded by other objects: **occlusion culling, backface culling**
- All work on the object level, instead of the primitive level



# CULLING



## VIEW FRUSTUM CULLING GOAL

Quickly eliminate large portions of the scene which will not be visible in the final image

- Not the exact visibility solution, but a quick-and-dirty conservative estimate of which primitives might be visible
  - Z-buffer & clipping this for the exact solution
- This conservative estimate is called the potentially visible set or PVS

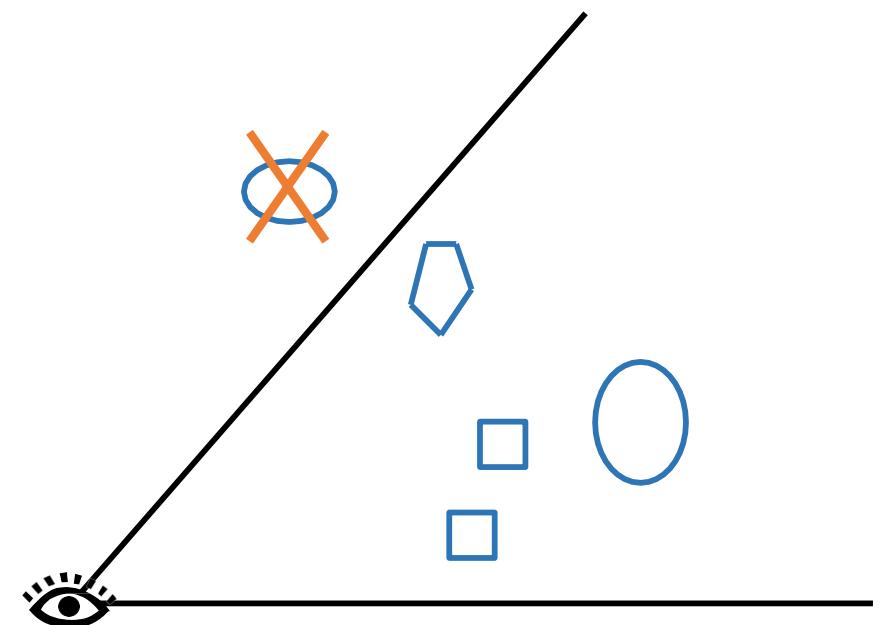


## VIEW FRUSTUM CULLING

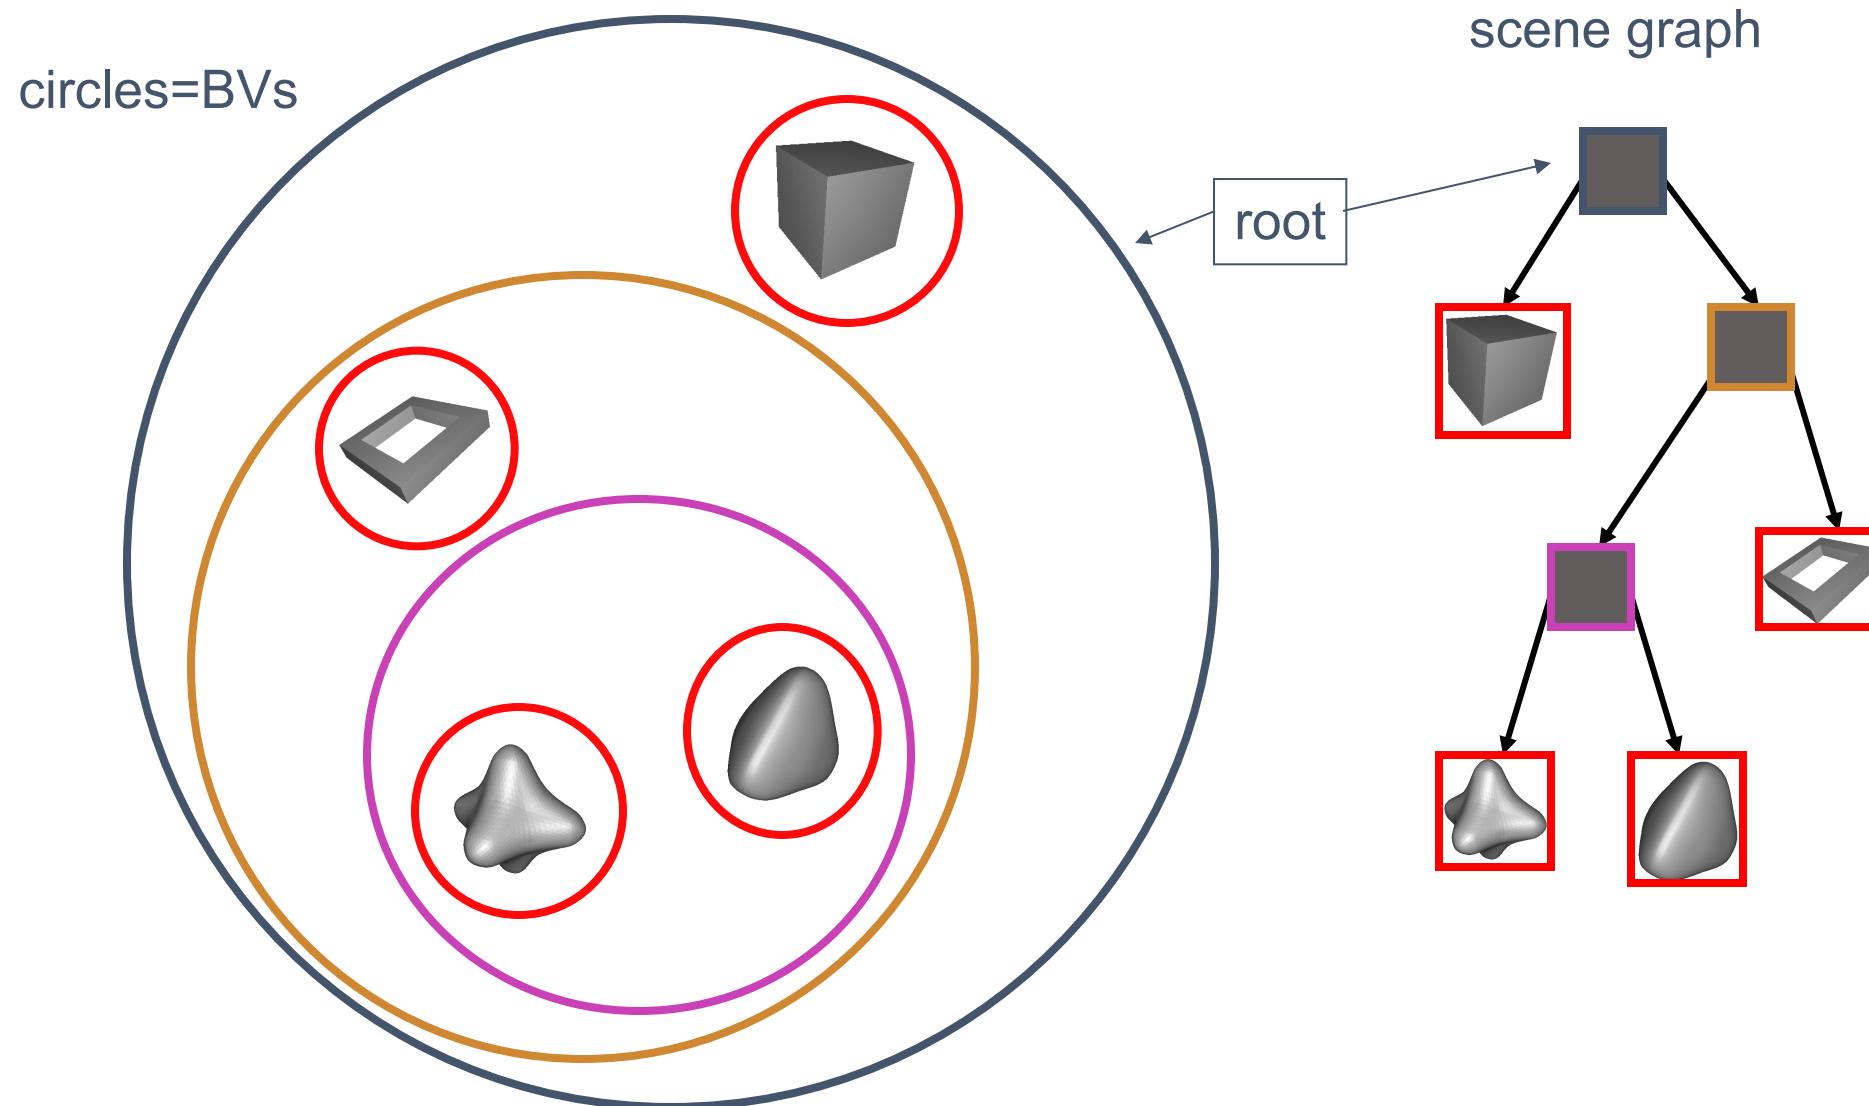
Remove object not inside view frustum (VF)

Complete test is called exclusion/inclusion/intersection test

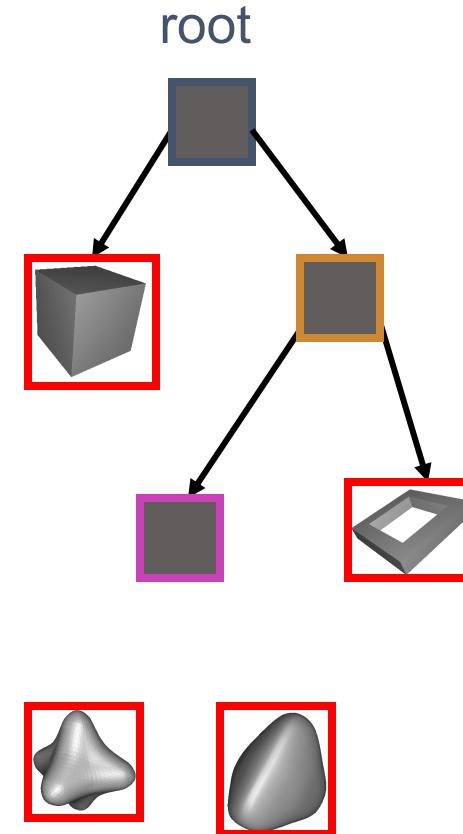
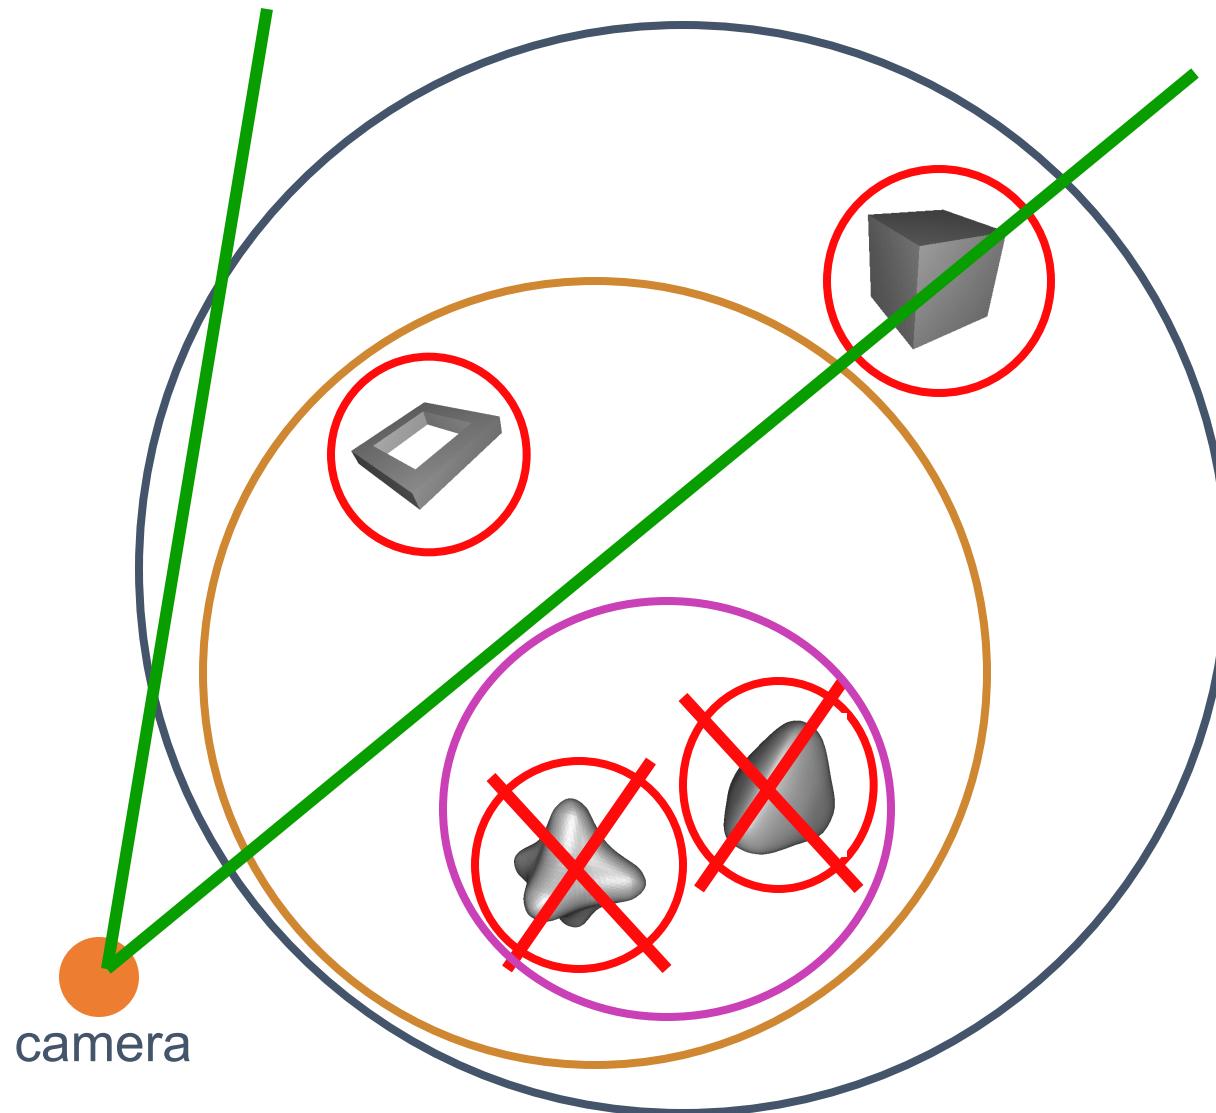
- Exclusion a BV is completely outside the VF, thus not rendered
- Inclusion the BV is completely inside the VF, thus rendered
- Intersection, part of the BV is inside, part is outside. Here we have a few options.



# SCENE GRAPH EXAMPLE



## EXAMPLE OF HIERARCHICAL VIEW FRUSTUM CULLING



# PORTAL CULLING

Goal: walk through architectural models (buildings, cities, catacombs)

Divide scene into cells

- Room, hallways
- Each object (furniture etc.) associated to a cell are referenced

Define portals

- Door/windows connect adjacent rooms
- For each cell associate portals to that cell

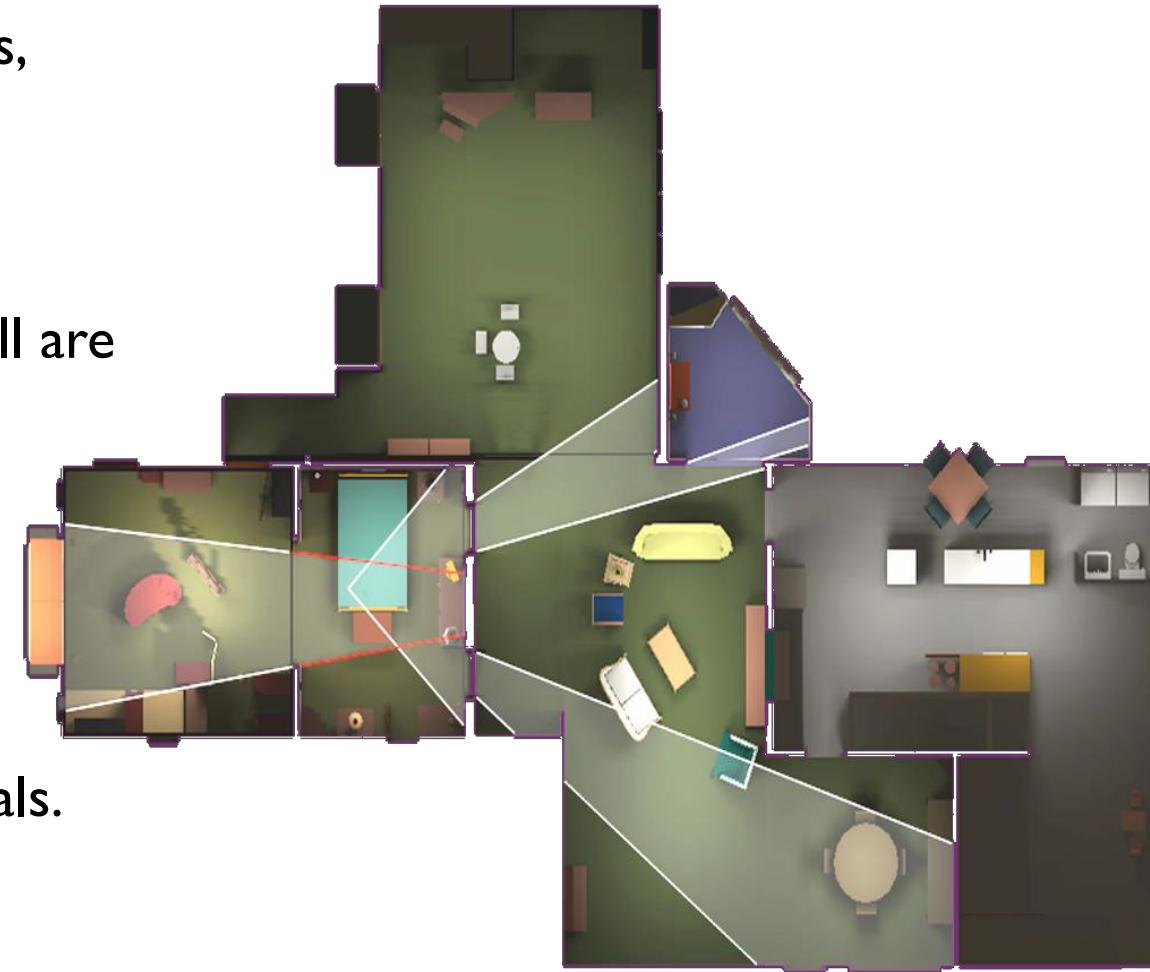
Create an adjacency graph of cells

- Each cell connect to other cells through portals.

Extension of VF culling

- The VF is adjusted to the adjacent portals

Notice: cells only see other cells through portals



## CELLS & PORTALS

Starting with cell containing eye-point, traverse graph, rendering visible cells using original View Frustum (VF)

For each Portal visible in VF

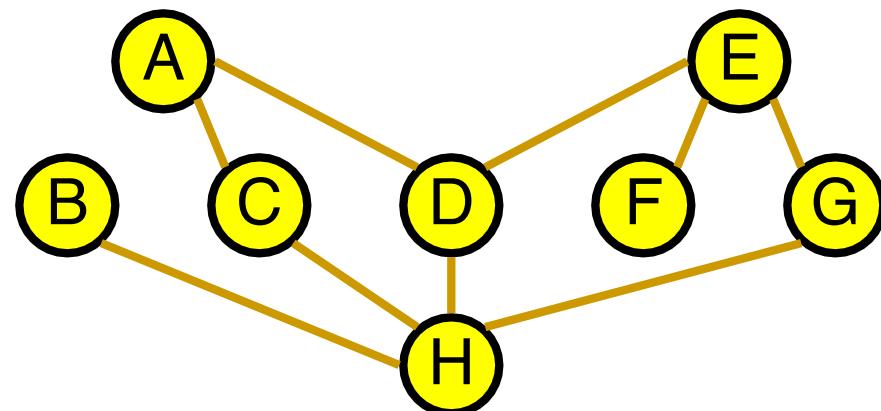
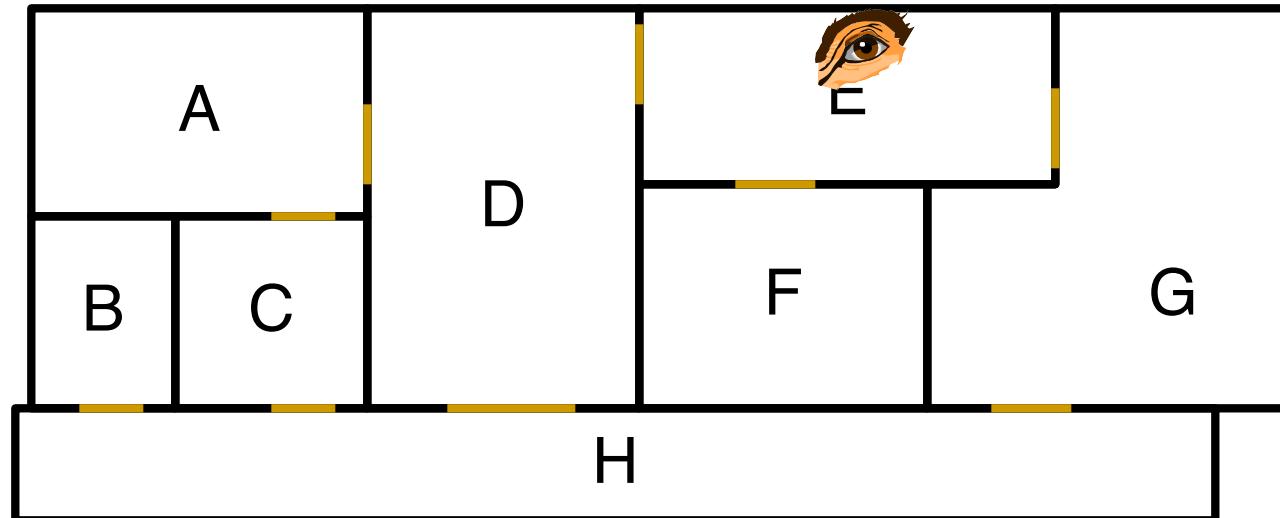
- Reduce VF to fit Portal
- Now render Cell with reduced VF
  - For each portal visible in VF... (recursion)

A cell is only visible if it can be seen through a sequence of portals

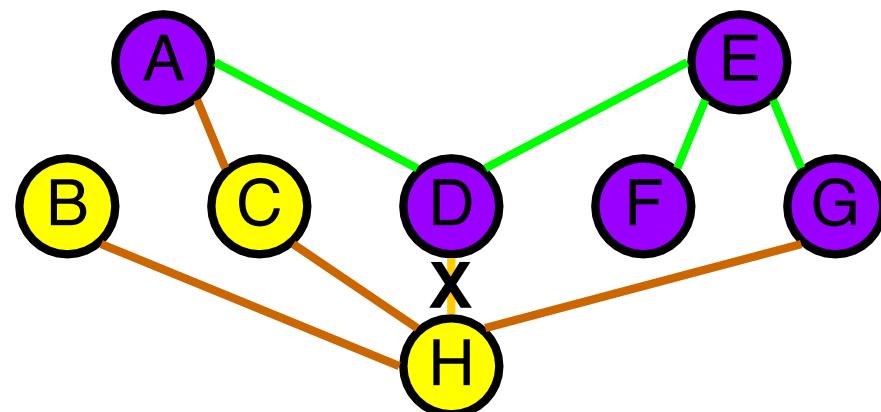
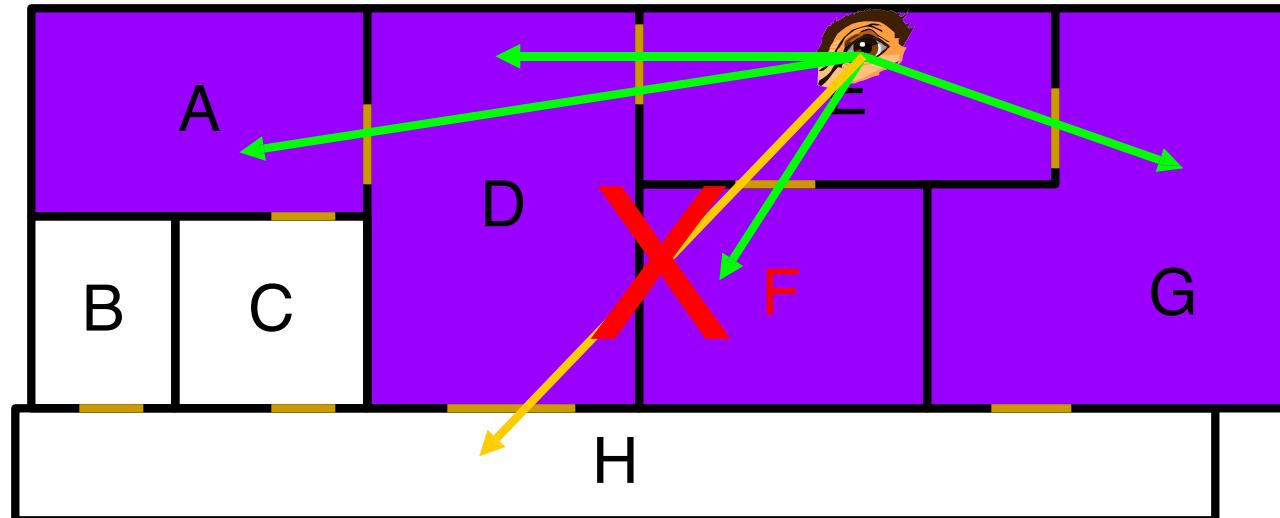
- So cell visibility reduces to testing portal sequences for a line of sight...



# CELLS & PORTALS



# CELLS & PORTALS



## CONTRIBUTION CULLING

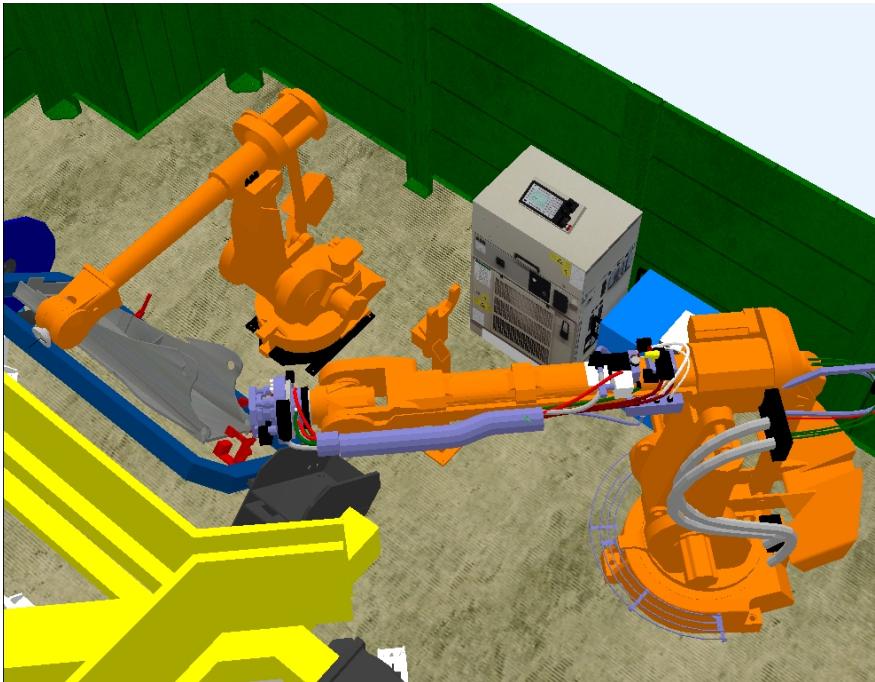
Idea: objects whose projected bounding volume occupy less than  $N$  pixels are culled

This is an approximative algorithm as the things you cull away may actually contribute to the final image

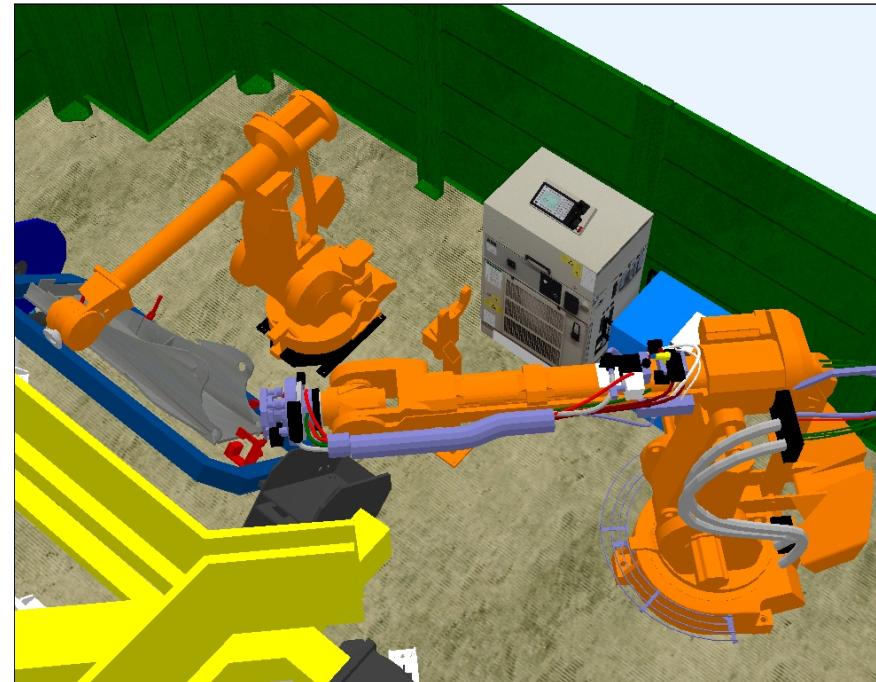
Advantage: trade-off quality/speed



# EXAMPLE OF CONTRIBUTION CULLING



**contribution culling OFF**



**contribution culling ON**

Not much visual difference, but 80-400% faster

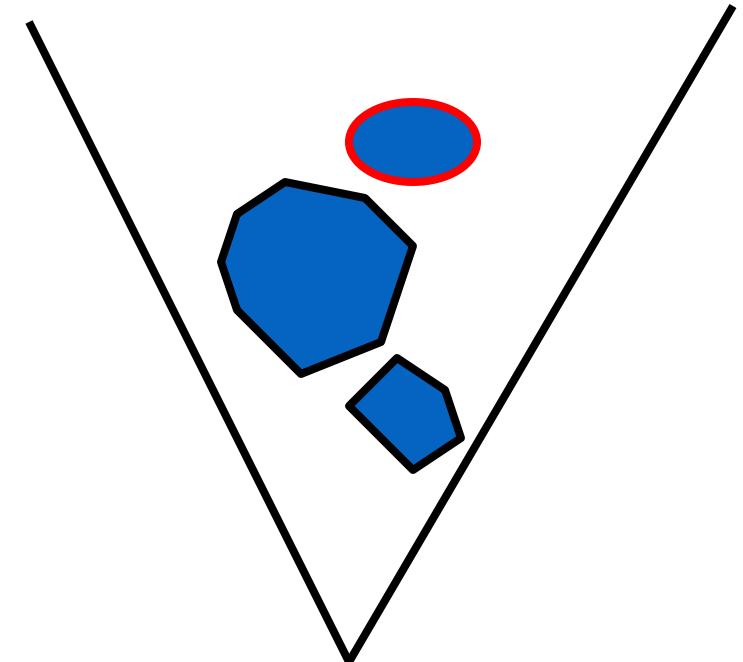
Particularly effective when moving



## OCCULTION CULLING

Main idea: Objects that lies completely “behind” another set of objects can be culled

Hard problem to solve efficiently

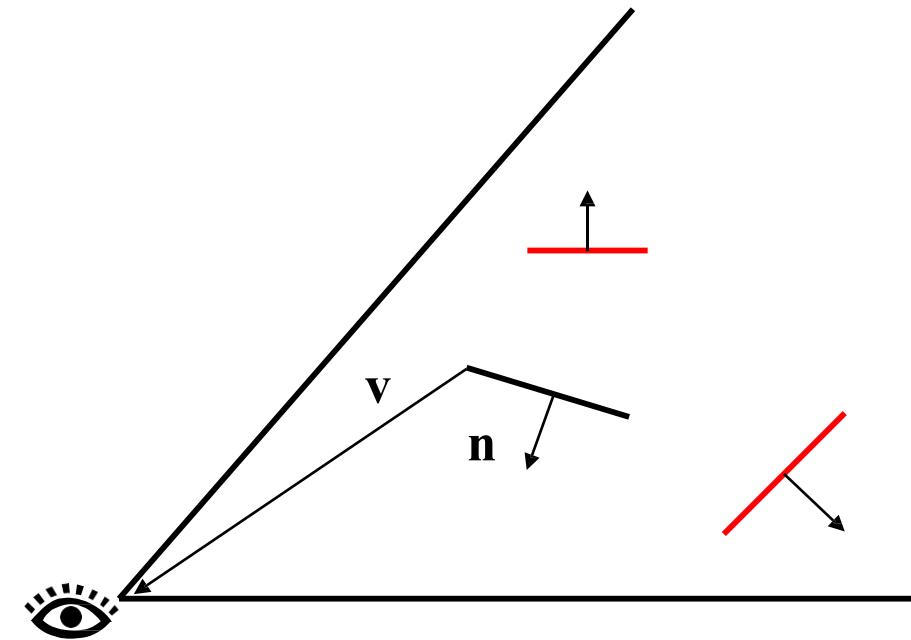


## BACKFACE CULLING

If normal of a surface is pointing away from viewer, remove it

Works for closed object where we don't want to see the inside of the object (for example a sphere)

Does not work well with for terrain (most polygons visible)



## SIMPLIFY GEOMETRY

One polygon

- Sprites
- Billboards
- Impostors

Many polygons

- Level of detail (LOD)



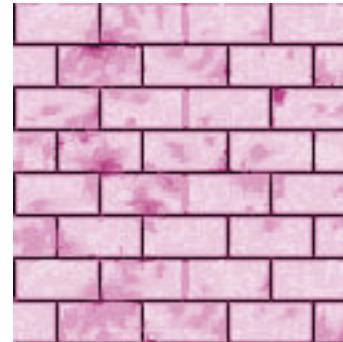
## OLD SCHOOL: SPRITES

A rectangular shaped image that moves around on the screen

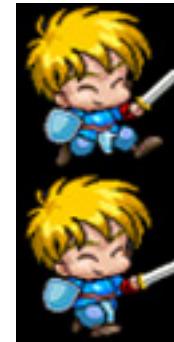
- One-to-one mapping with pixels on the screen
- The earliest PC graphics were sprite

### Implementation

- An image texture on a polygon with the use of alpha channel



+



sprite image

=



result



# BILLBOARDING

Textured polygon

Combined with alpha  
texturing and (potentially)  
animation

- ex. smoke, fire, explosions,  
vapor trails, clouds



## Axial Billboard

Billboards at fixed location  
rotates to face user

Useful for representing  
objects with cylindrical  
symmetry

- displaying trees

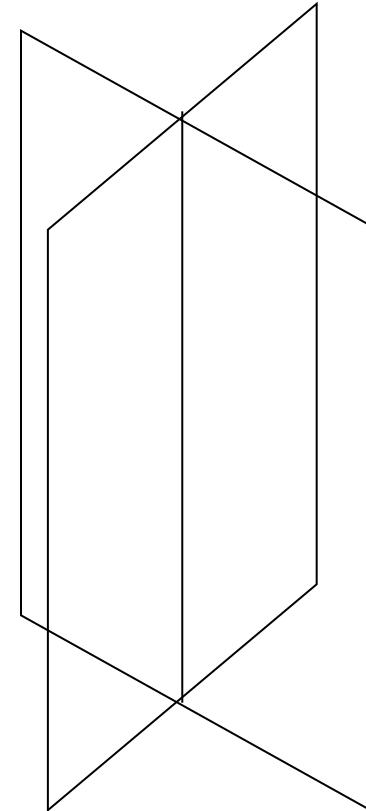


## MULTI-POLYGON BILLBOARDS

Use two polygons at right angles

- No alignment with viewer

Use more polygons for better appearance



## IMPOSTORS

Replace 3D geometry with a 2D image

2D image fools viewer into thinking 3D  
geometry is still there

Big limitation: No parallax

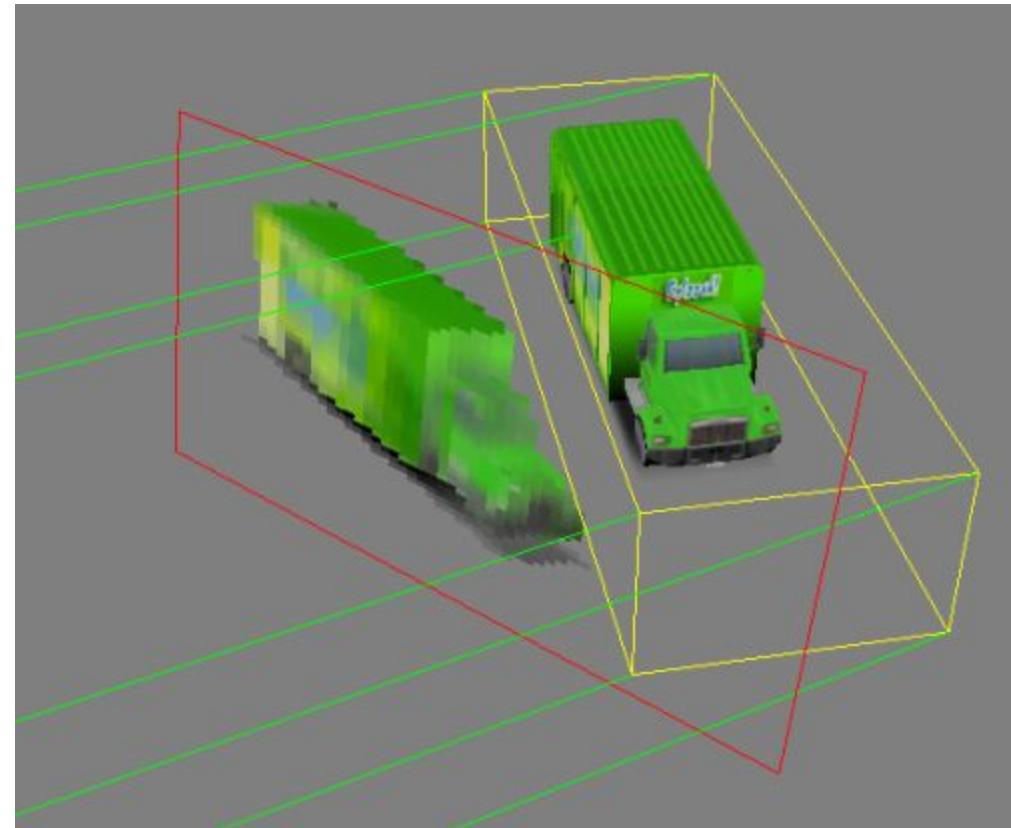


# IMPOSTORS TECHNIQUE

Basic idea

- Render set of geometry into a large texture: an impostor
- Now render impostor texture instead of the geometry

Only helpful when impostors can be reused across many frames



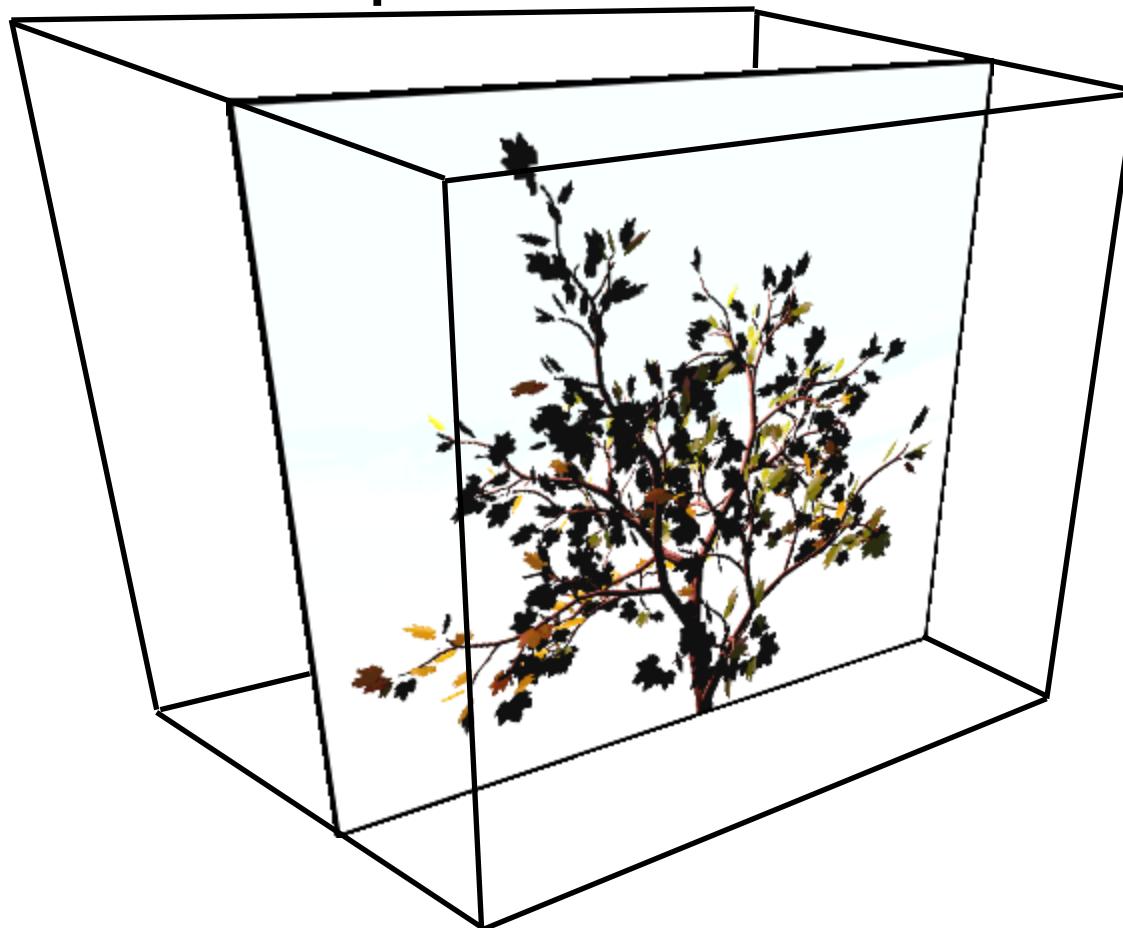
## IMPOSTORS: EXAMPLE

We render a set of geometry into an impostor (image/texture)



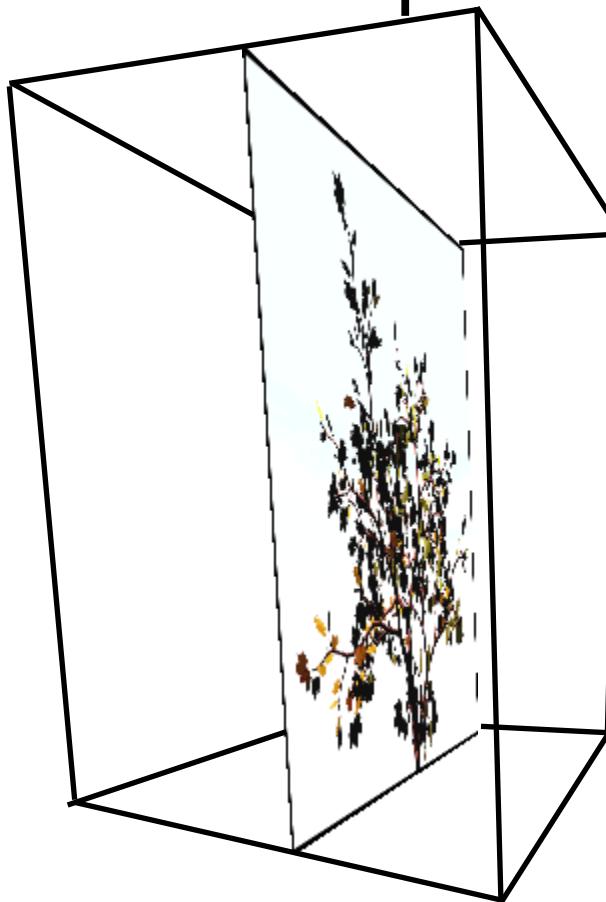
## IMPOSTORS: EXAMPLE

We can re-use this impostor in 3D for several frames



## IMPOSTORS: EXAMPLE

Eventually, we have to update the impostor



## LEVEL OF DETAIL

The problem:

- Geometric datasets can be too complex to render at interactive rates

One solution:

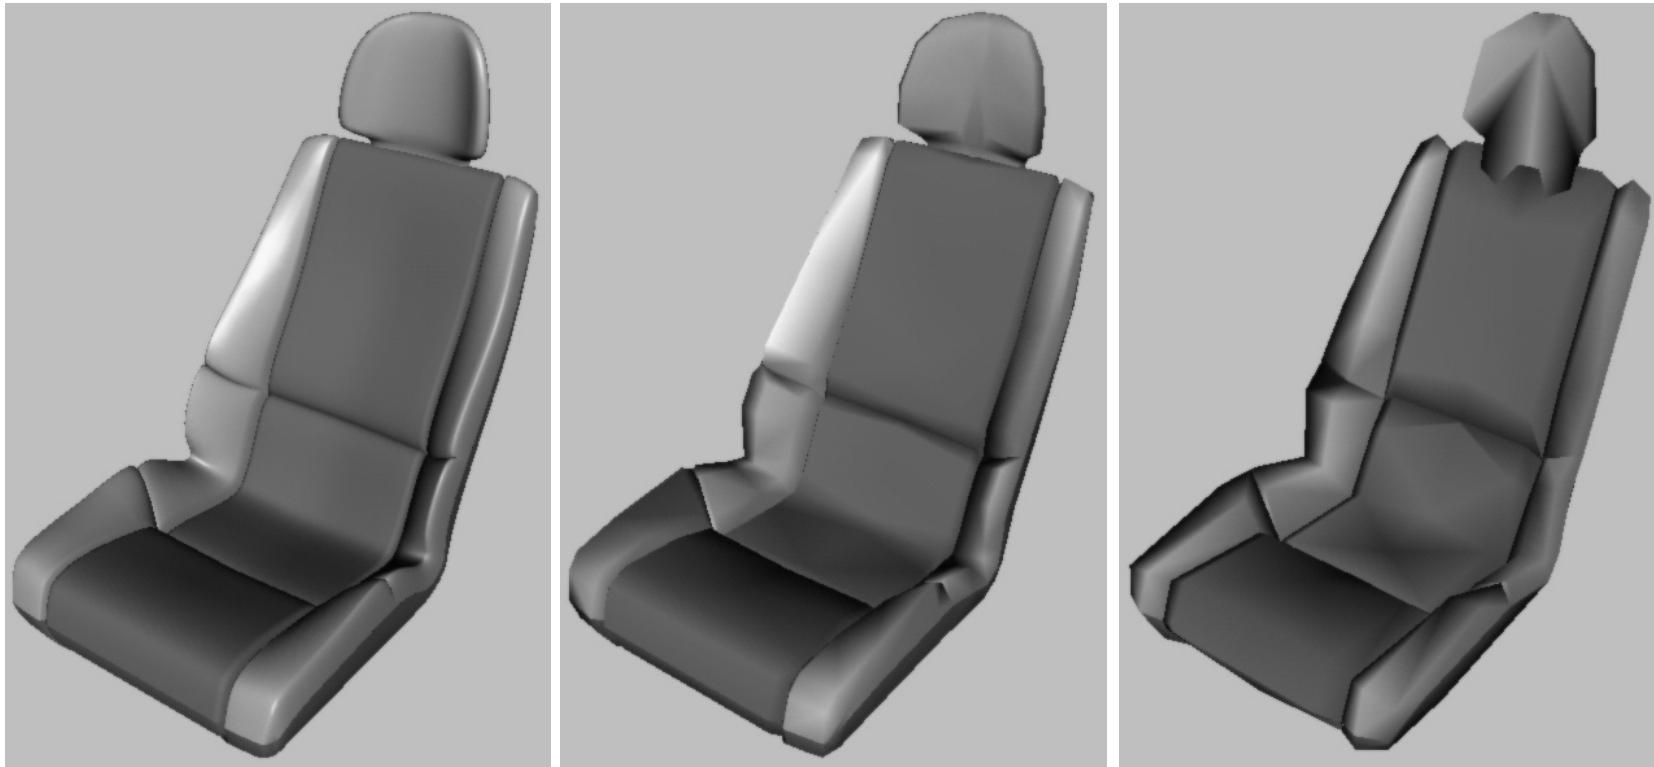
- Simplify the polygonal geometry of small or distant objects
- Known as ***Level of Detail*** or ***LOD***
  - A.k.a. polygonal simplification, geometric simplification, mesh reduction, decimation, multiresolution modeling, ...
- Focuses on the fidelity / performance tradeoff



## LEVEL-OF-DETAIL RENDERING

Use different levels of detail at different distances from the viewer

More triangles closer to the viewer



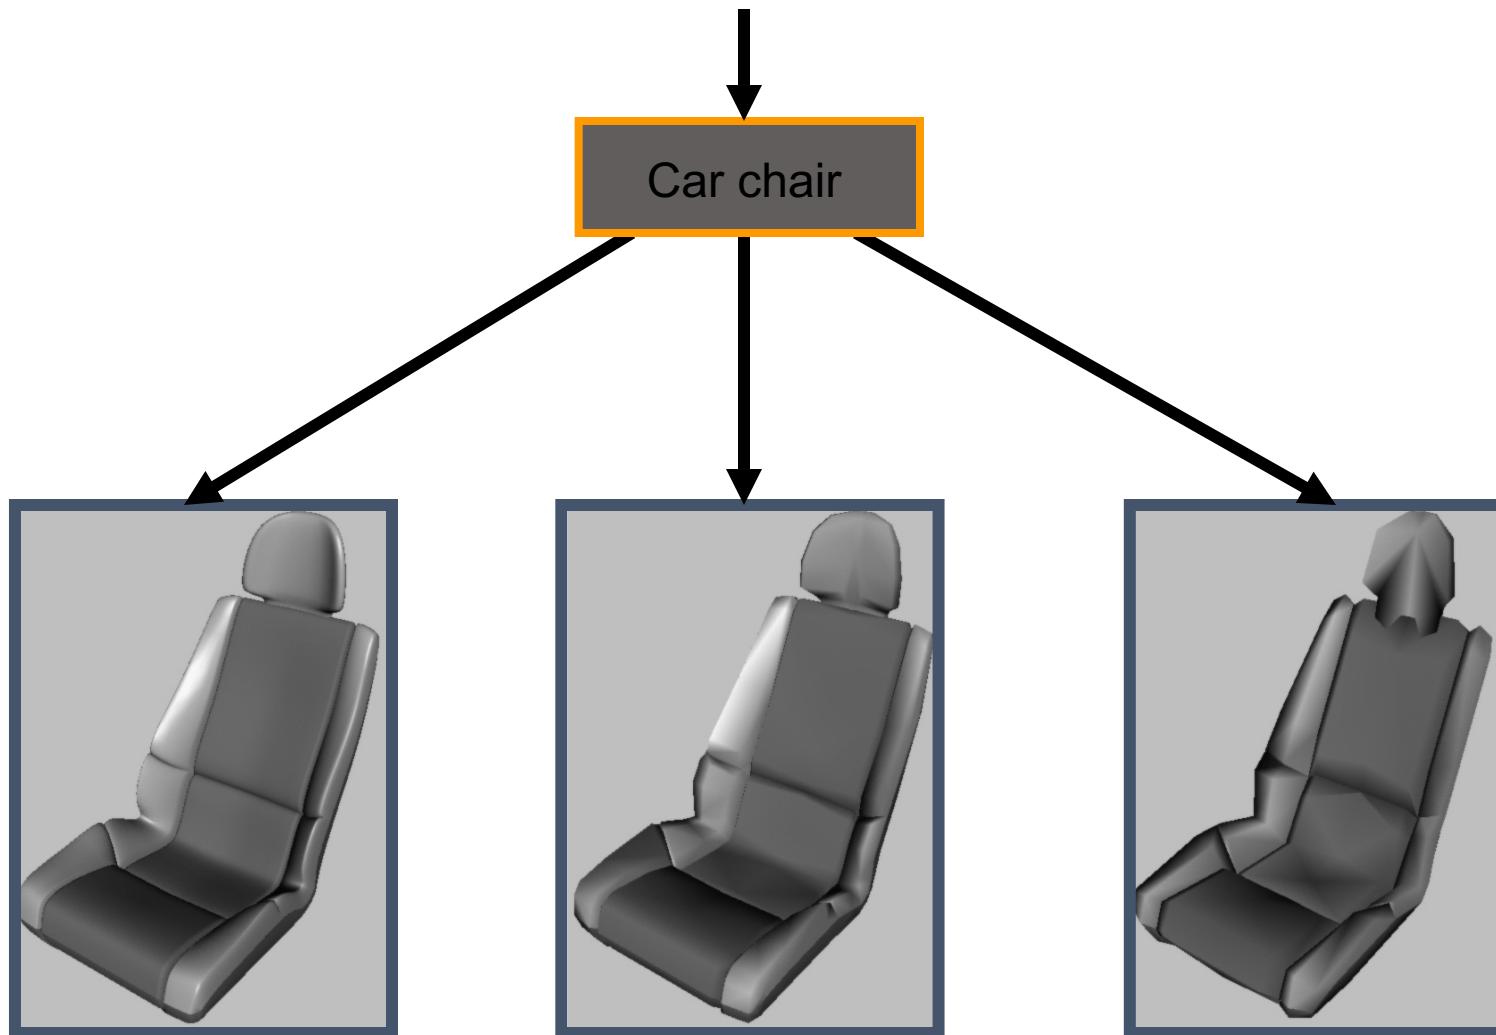
## LOD RENDERING

Use area of projection of BV to select appropriate LOD

Not much visual difference, but a lot faster



# SCENE GRAPH WITH LODS



## ANOTHER EXAMPLE



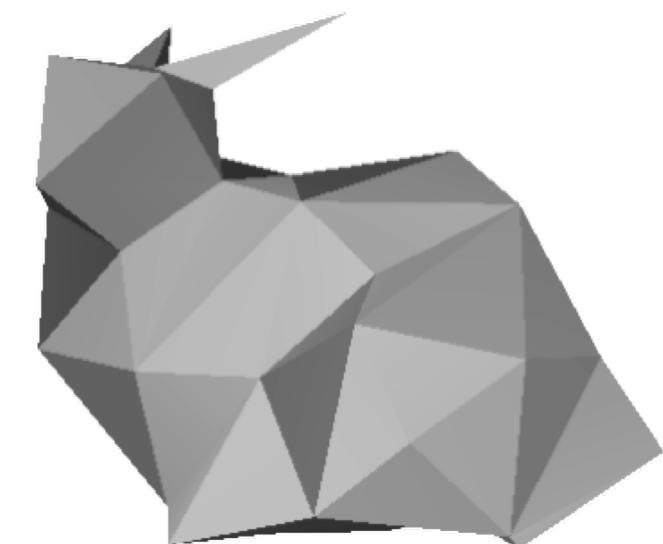
69,451 polys



2,502 polys



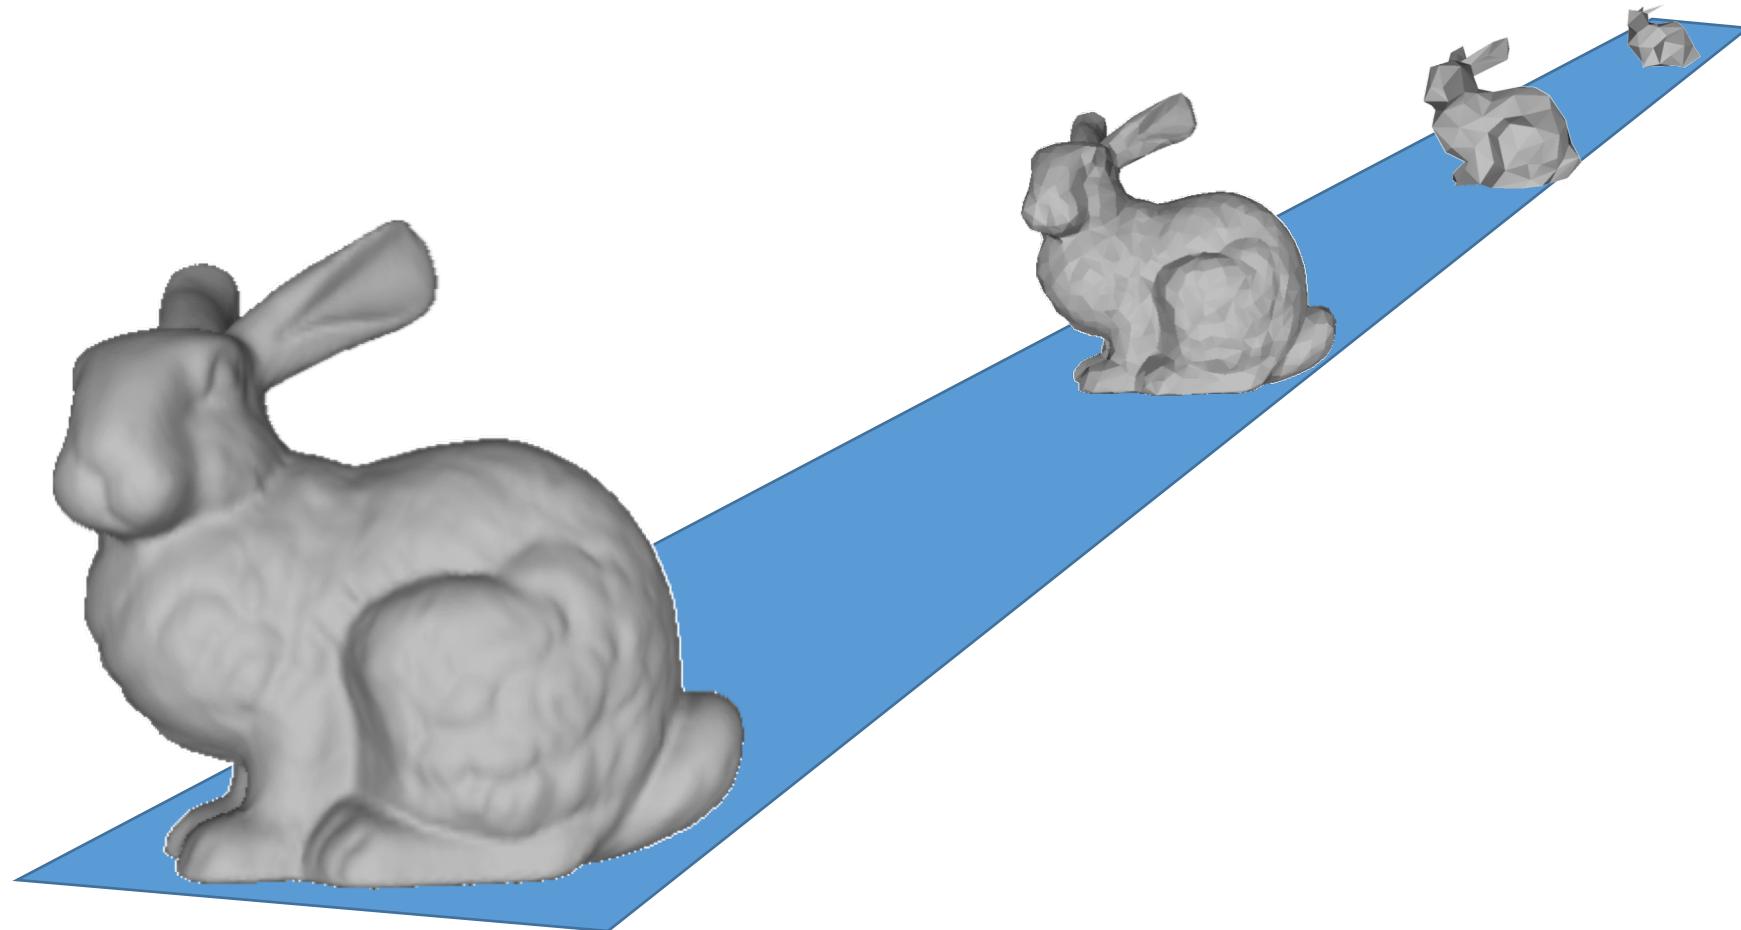
251 polys



76 polys



# ANOTHER EXAMPLE



COURTESY STANFORD 3D SCANNING REPOSITORY



## LEVEL OF DETAIL: THE BIG QUESTIONS

How to generate simpler versions of a complex model?

How to evaluate the fidelity of the simplified models?

When to use which LOD of an object?



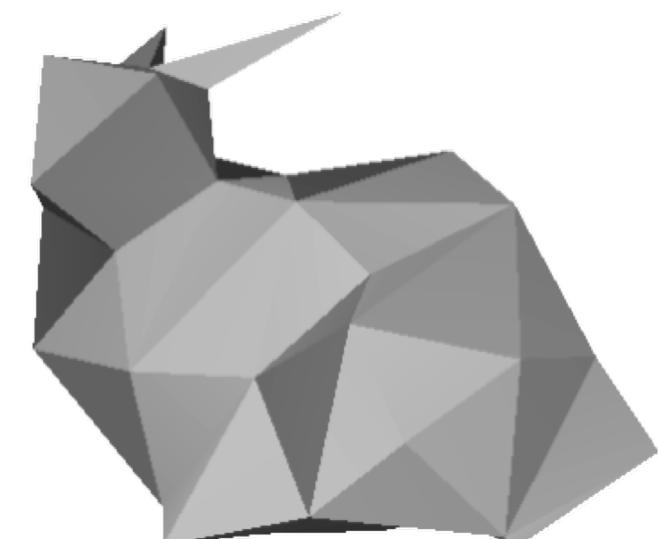
69,451 polys



2,502 polys



251 polys



76 polys



## DISCRETE LOD APPROACH

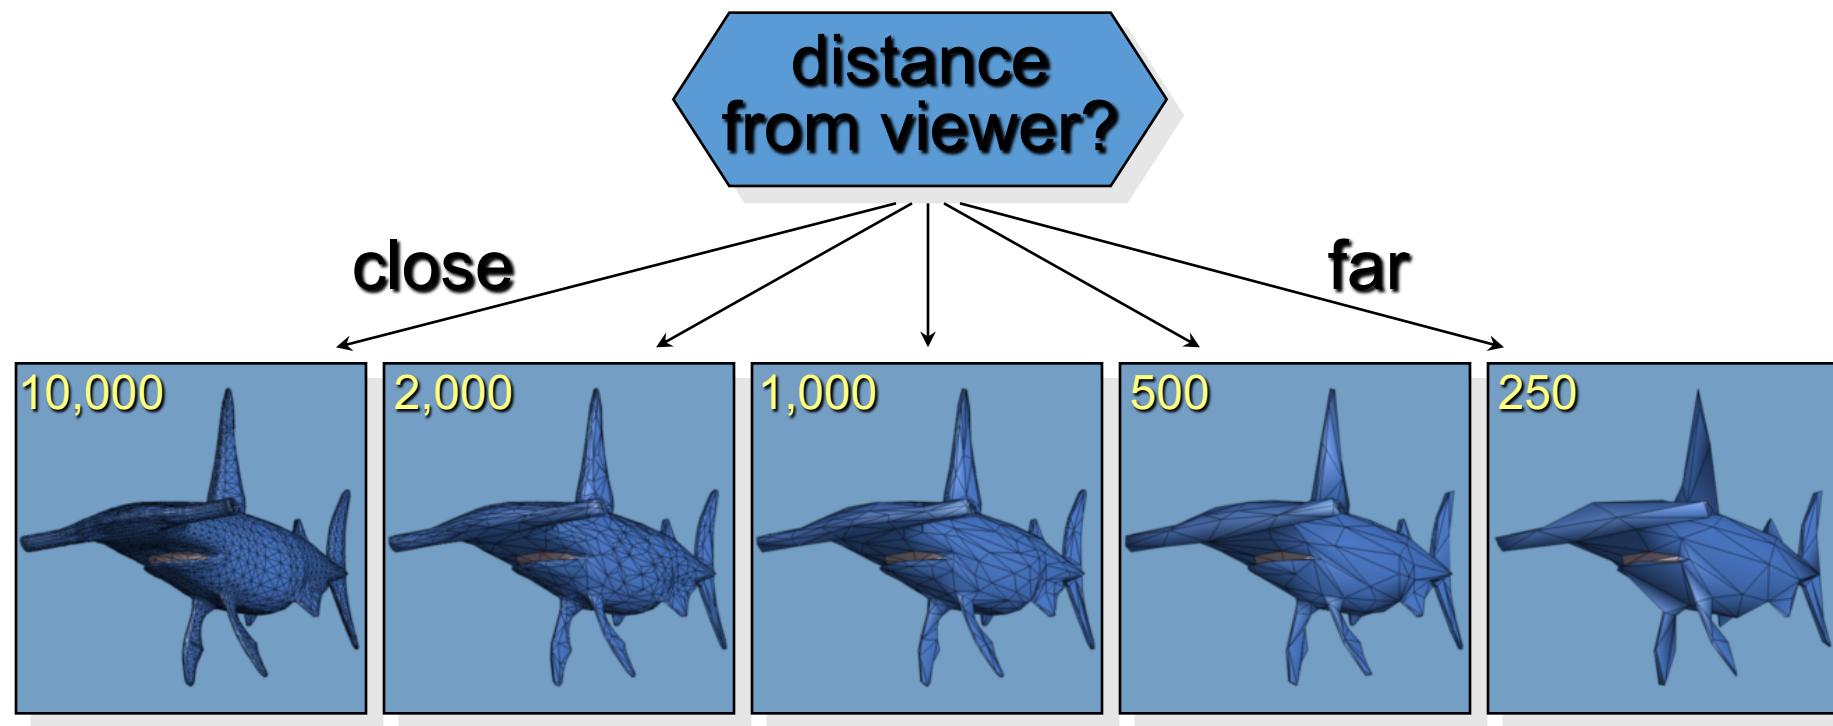
Traditional LOD in a nutshell:

- Create LODs for each object separately in a preprocess
- At run-time, pick each object's LOD according to the object's distance (or similar criterion)

Since LODs are created at fixed resolutions, we call this discrete LOD



# DISCRETE LEVEL-OF-DETAIL (LOD)



## DISCRETE LOD:ADVANTAGES

Simple programming model—decouples simplification and rendering

- LOD creation need not address real-time rendering constraints
- Run-time rendering need only pick LODs



Level 0

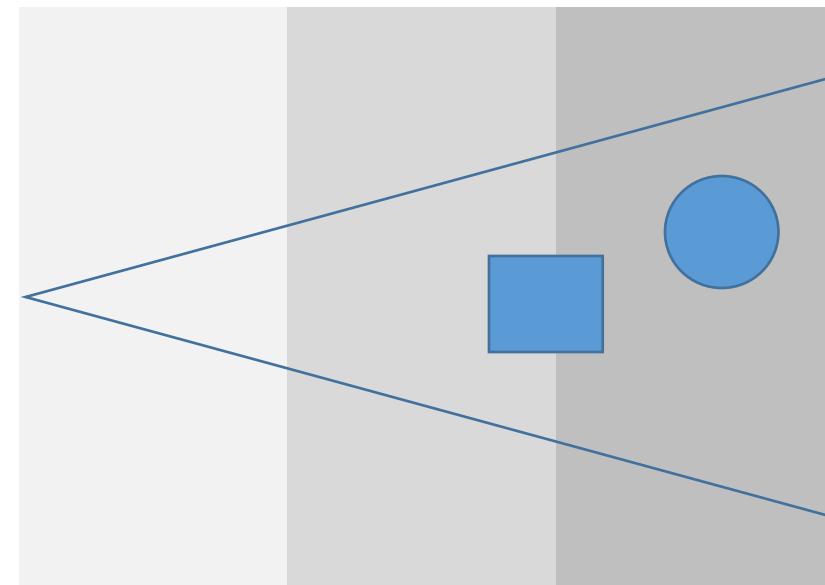
Level 1

Level 2

## CHOOSING THE LOD

### Heuristic Approach

- Assign each LOD a range of distances
- Calculate distance from viewer to object
- Use corresponding LOD



### Problems:

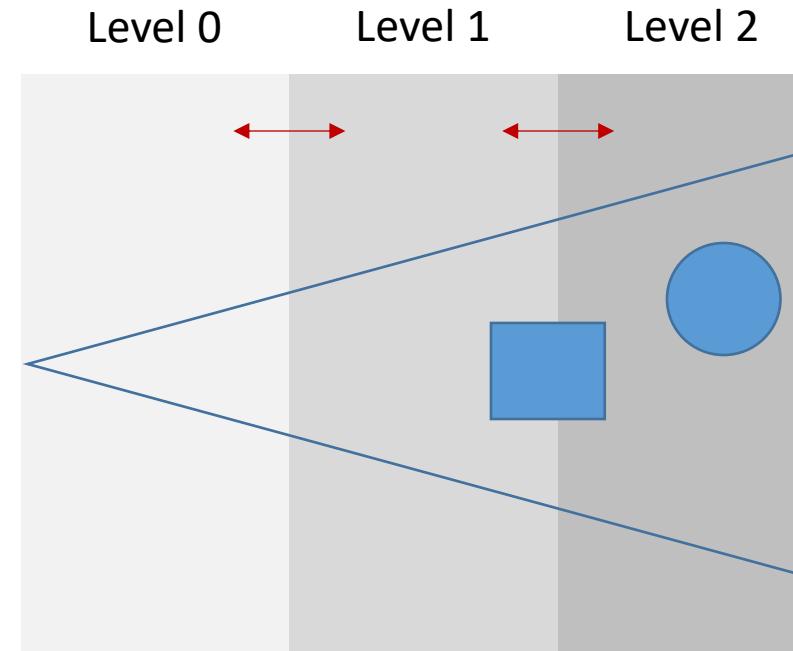
- Visual “pop” when switching LODs can be disconcerting
- Doesn’t maintain constant frame rate
- Requires someone to assign switching distances by hand
- Correct switching distance may vary with field of view, resolution, etc.



# CHOOSING THE LOD

Adaptively maintaining frame rate

- Scale LOD switching distances by a “bias”
  - If last frame took too long, decrease bias
  - If last frame took too little time, increase bias



Problems:

- Oscillation caused by overly aggressive feedback
- Sudden change in rendering load can still cause overly long frame times

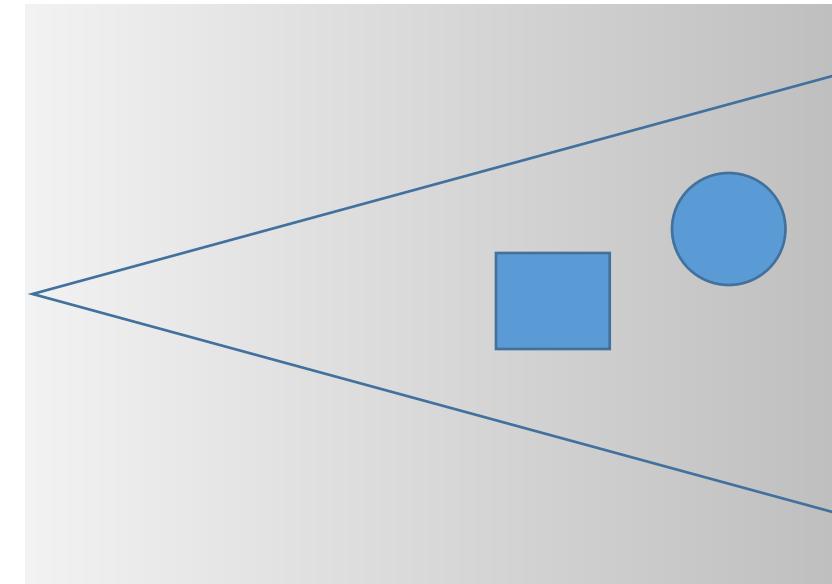


# CHOOSING THE LOD

Predictively maintaining frame rate

For each LOD estimate:

- Cost (rendering time)
  - # of polygons
  - Vertex & fragment processing load
- Benefit (importance to the image)
  - Size: larger objects contribute more to image
  - Accuracy: # of verts/polys, shading model, etc.
  - Priority: account for inherent importance
  - Eccentricity: peripheral objects harder to see
  - Velocity: fast-moving objects harder to see
  - Hysteresis: avoid flicker; use previous frame state



Given a fixed time budget, select LODs to maximize benefit within a cost constraint  
[Funkhouser & Sequin, SIGGRAPH 93]

- Sort objects by benefit/cost ratio, pick in sorted order until budget is exceeded



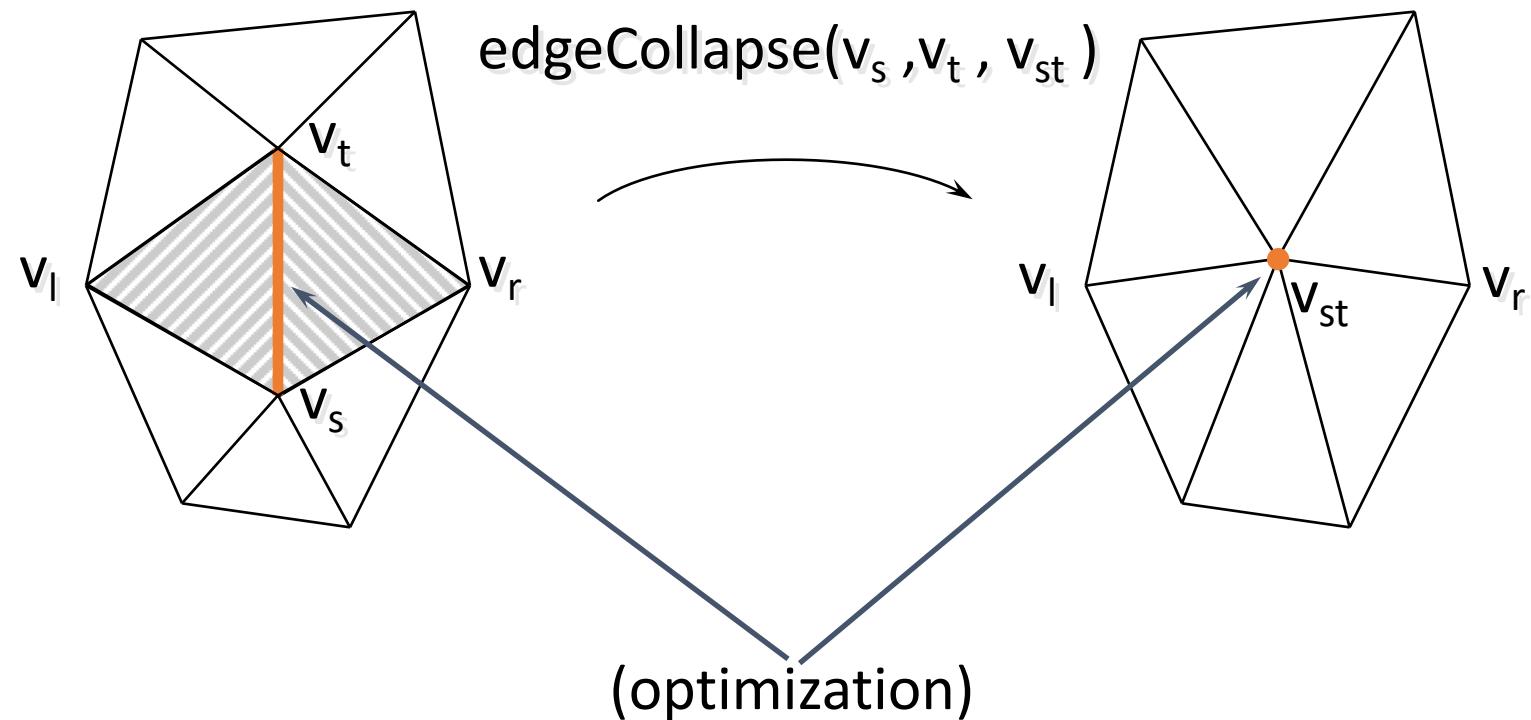
## CONTINUOUS LEVEL OF DETAIL

Create data structure from  
which a desired level of detail  
can be extracted *at run time.*

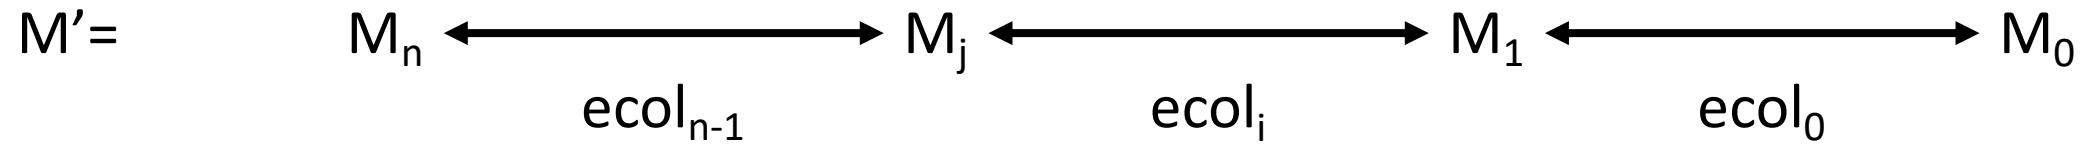
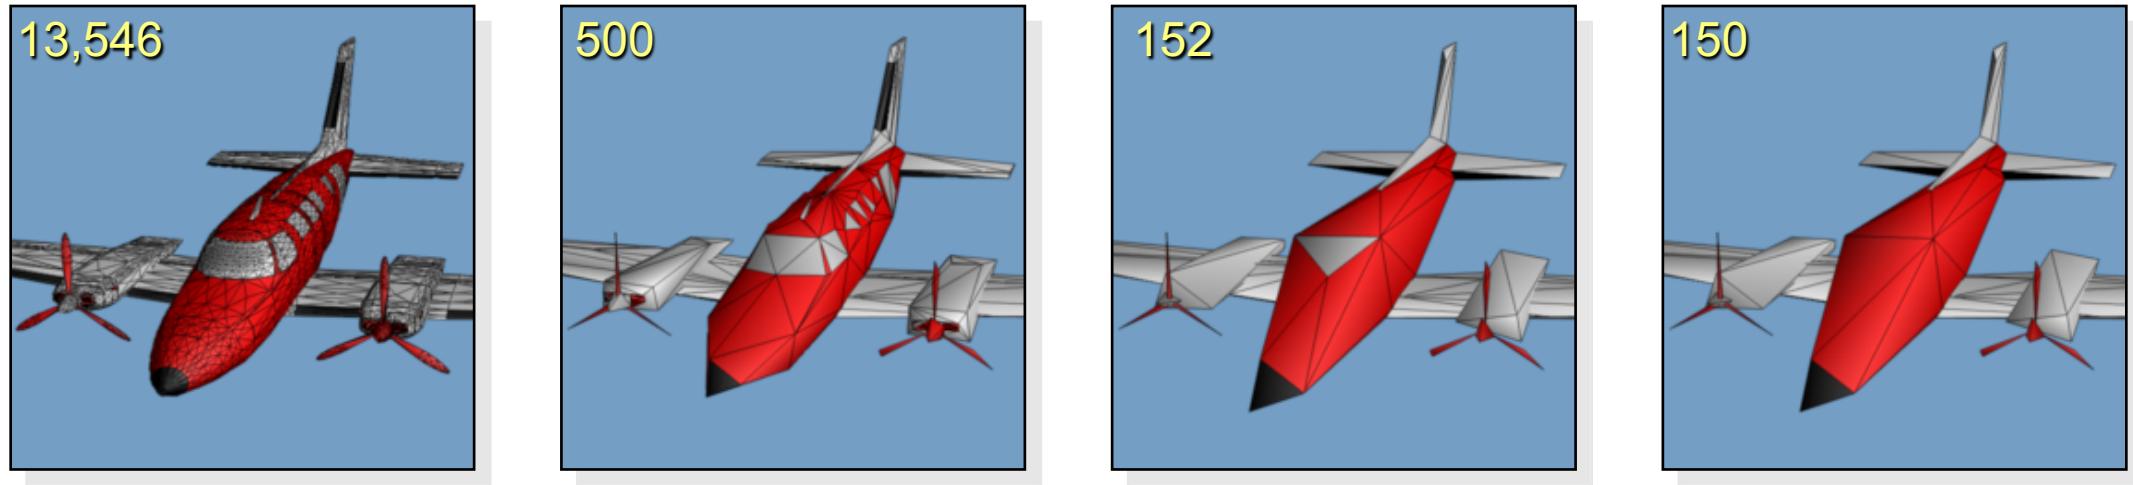


# NEW MESH SIMPLIFICATION PROCEDURE

Apply sequence of edge collapses



# SIMPLIFICATION PROCESS (ALSO INVERTIBLE)



## CONTINUOUS LOD:ADVANTAGES

Better granularity == better fidelity

- LOD is specified exactly, not chosen from a few pre-created options
- Thus objects use no more polygons than necessary, which frees up polygons for other objects
- Net result: better resource utilization, leading to better overall fidelity/polygon

Better granularity == smoother transitions

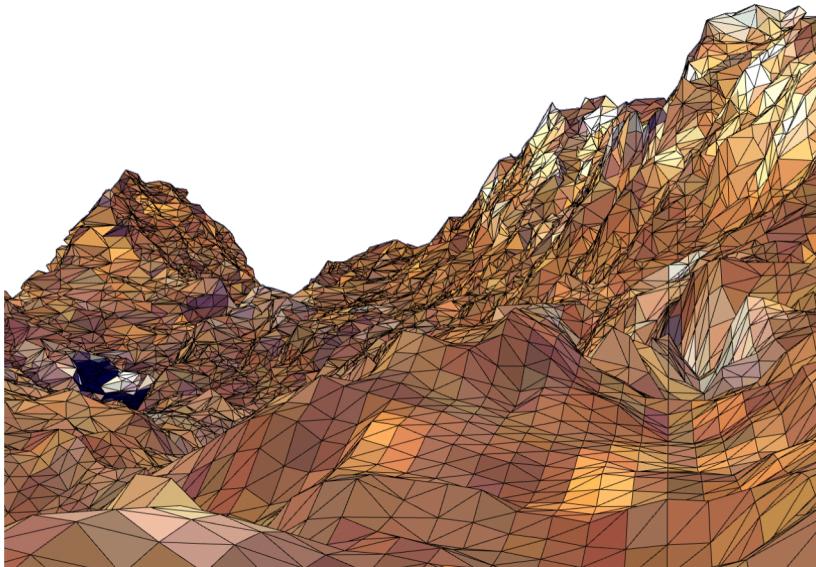
- Switching between traditional LODs can introduce visual “popping” effect
- Continuous LOD can adjust detail gradually and incrementally, reducing visual pops
  - Can even geomorph the fine-grained simplification operations over several frames to eliminate pops [Hoppe 96, 98]

Naturally leads to view-dependent LOD

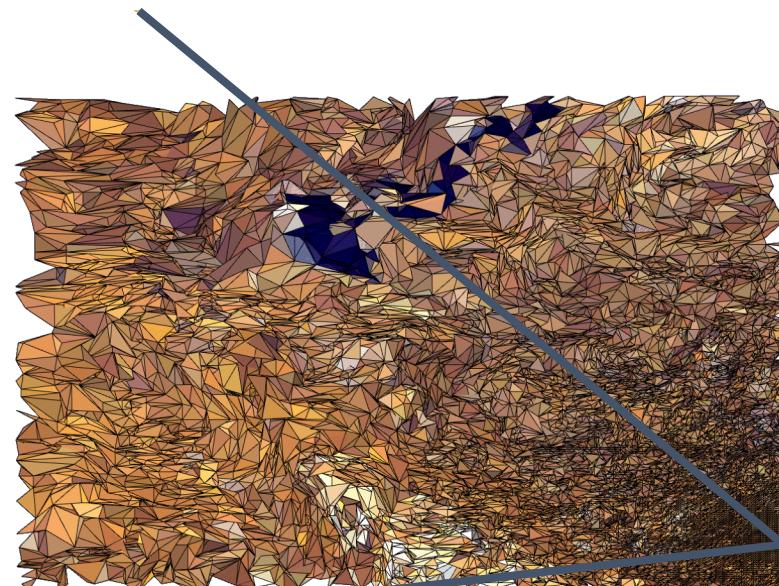


## VIEW-DEPENDENT LOD: EXAMPLES

Use current view parameters to select best representation *for the current view*  
Show nearby portions of object at higher resolution than distant portions



View from eyepoint



Birds-eye view



# DEMO VIDEO

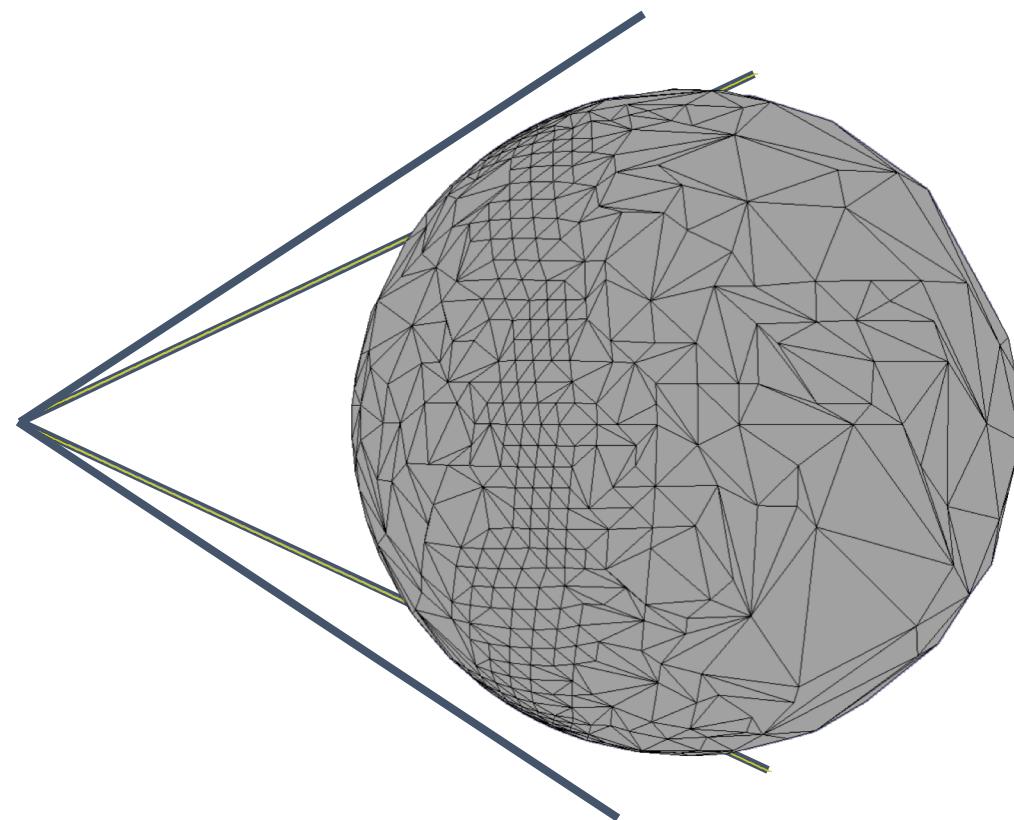


[HTTPS://YOUTU.BE/JBVYQU\\_wFEE](https://youtu.be/JBVYQU_wFEE)



## VIEW-DEPENDENT LOD: EXAMPLES

Show silhouette regions of object at higher resolution than interior regions



## VIEW-DEPENDENT LOD

### Advantages

- Enables drastic simplification of very large objects
- Better granularity
  - Allocates polygons where they are most needed
  - Enables even better overall fidelity

### Disadvantages

- More complicated object data structures
- More expensive LOD selection



